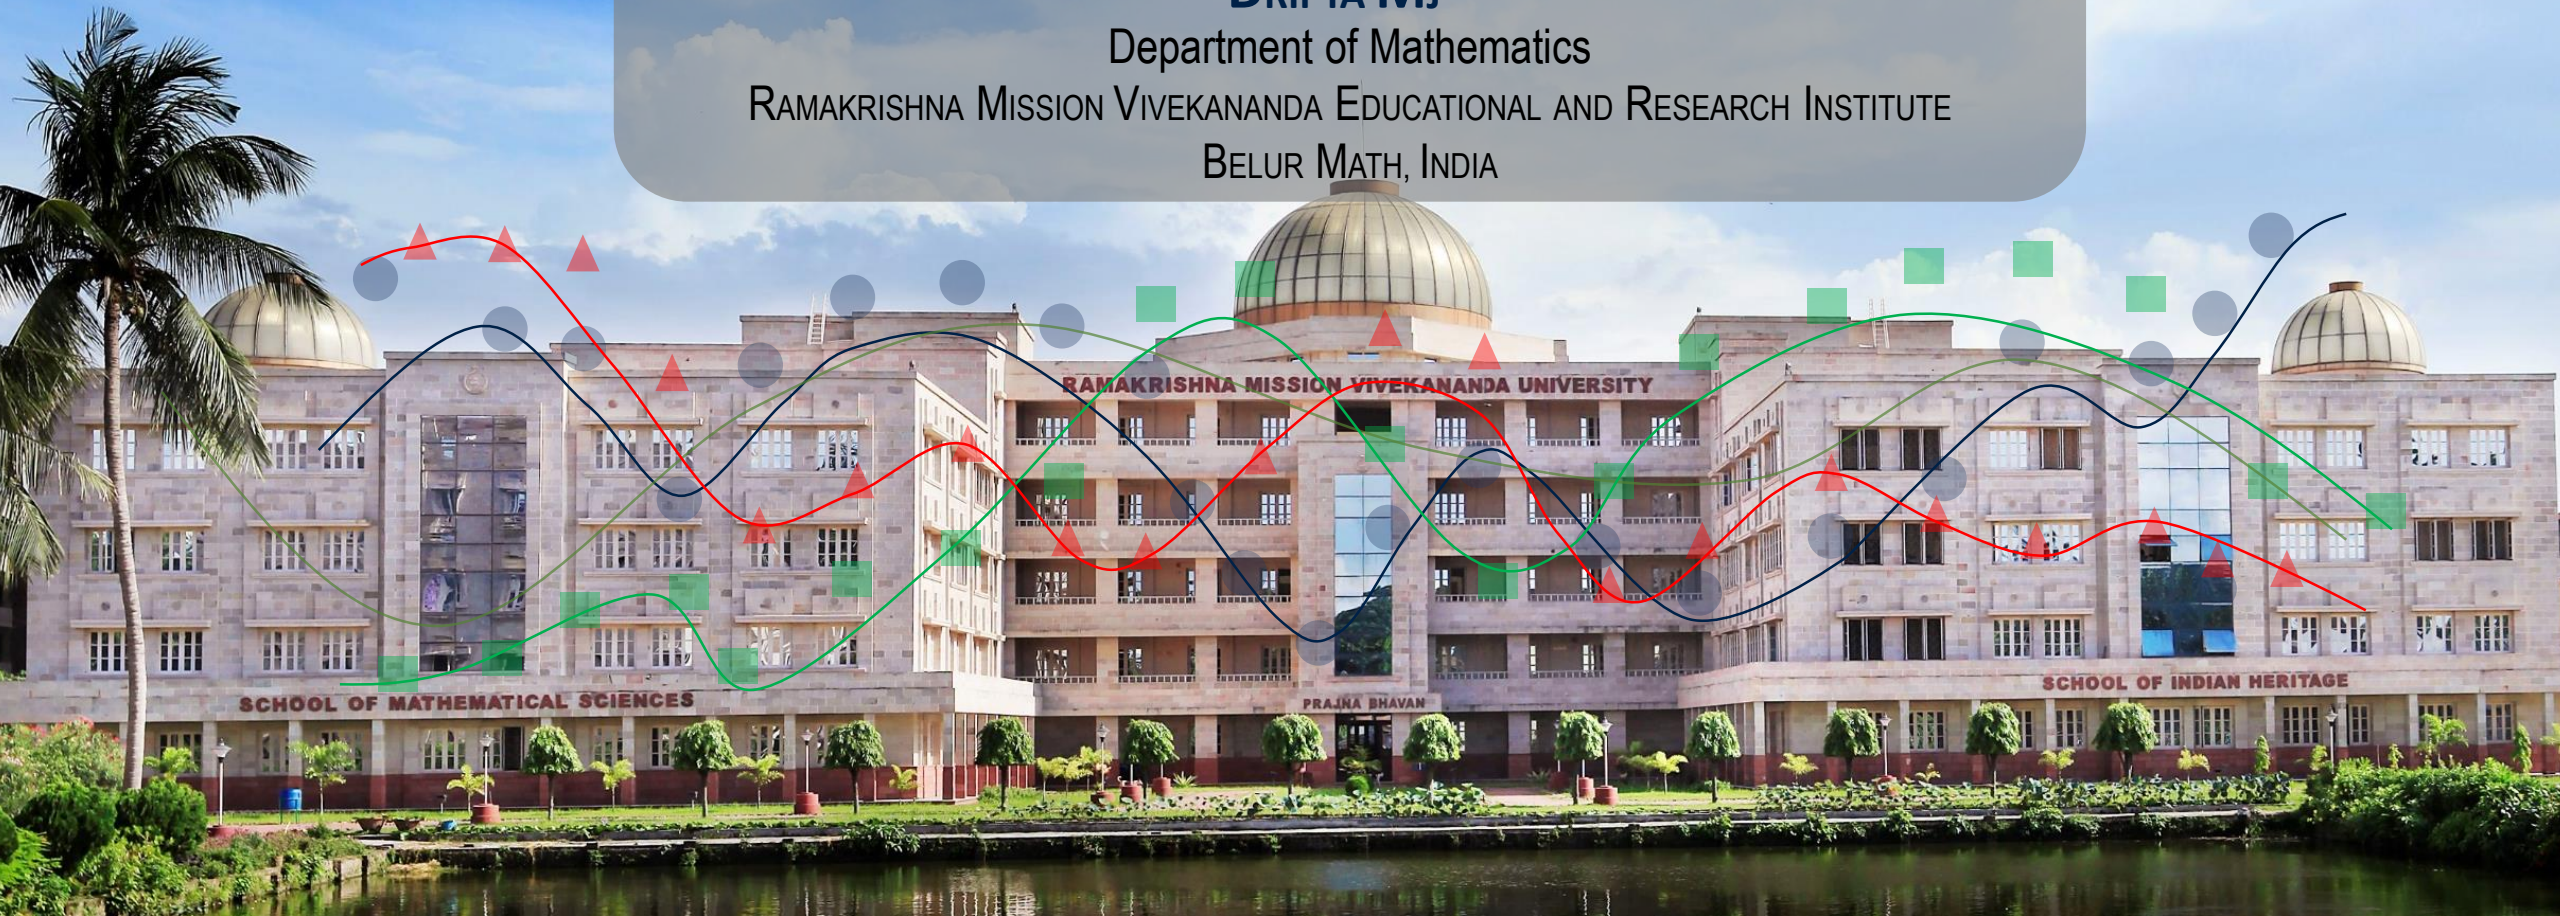


# Attention Mechanism

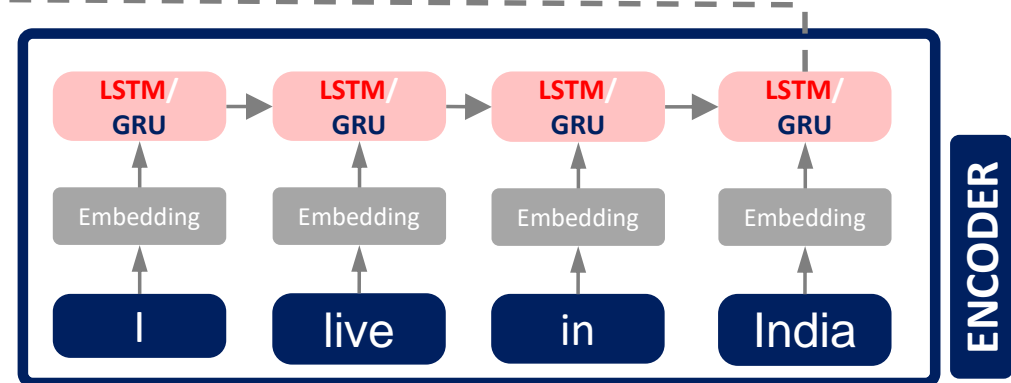
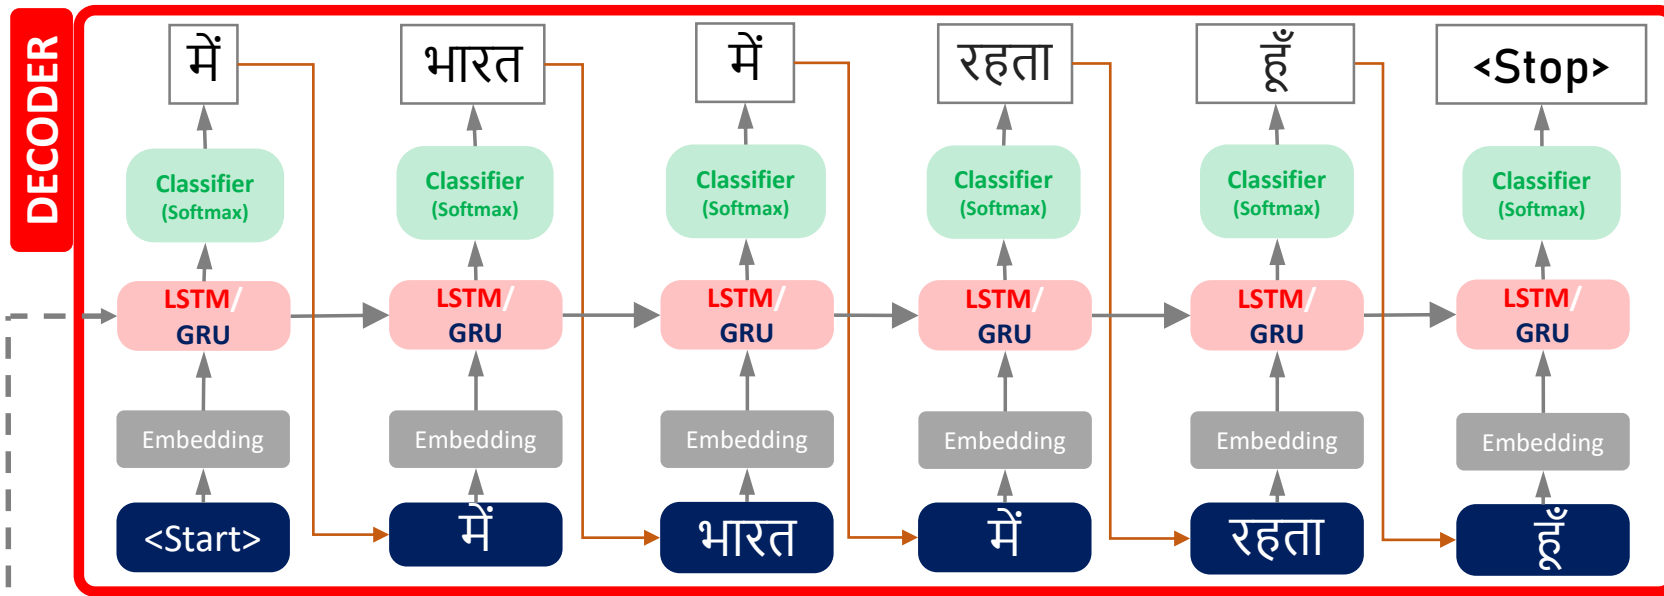
**DRIPTA MJ**

Department of Mathematics

RAMAKRISHNA MISSION VIVEKANANDA EDUCATIONAL AND RESEARCH INSTITUTE  
BELUR MATH, INDIA

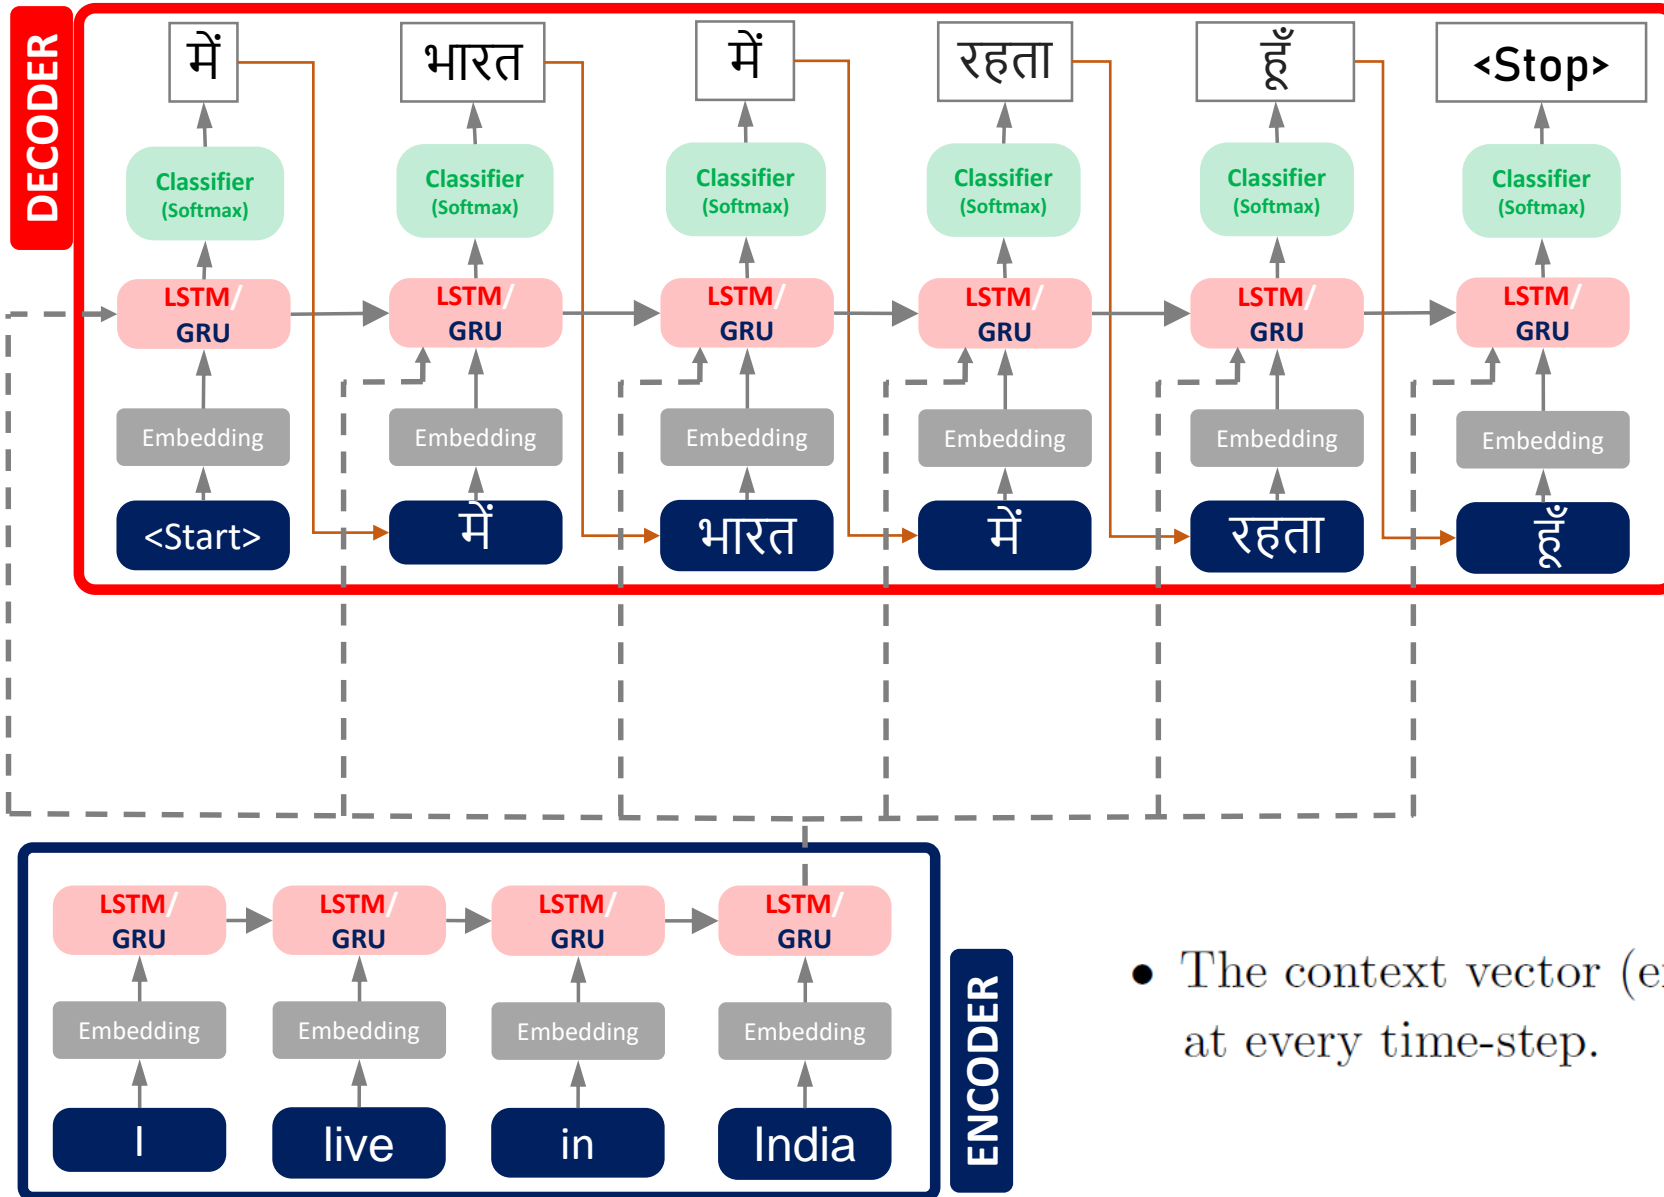


# Translation: Review



- The context vector (encoding) can be passed to the decoder as initialization of the hidden state (of decoder).

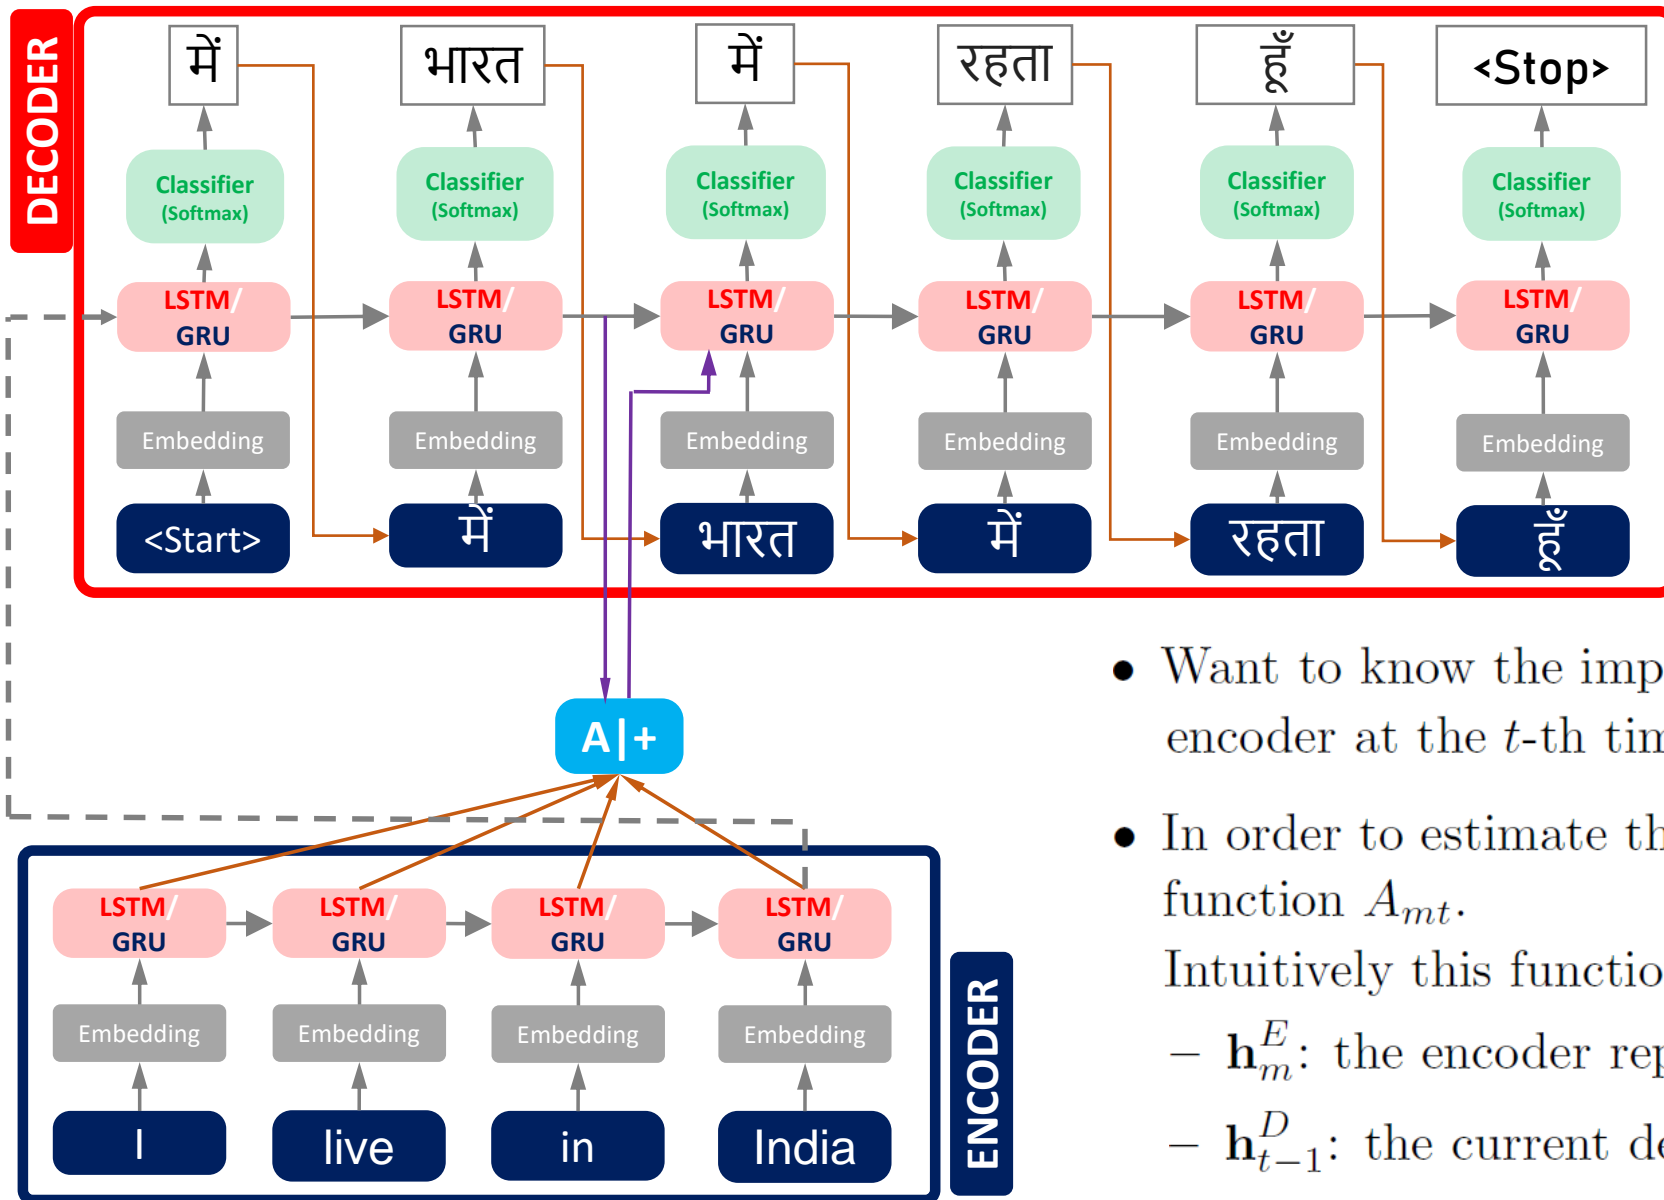
# Translation: Review



- The context vector (encoding) can be passed to the decoder at every time-step.

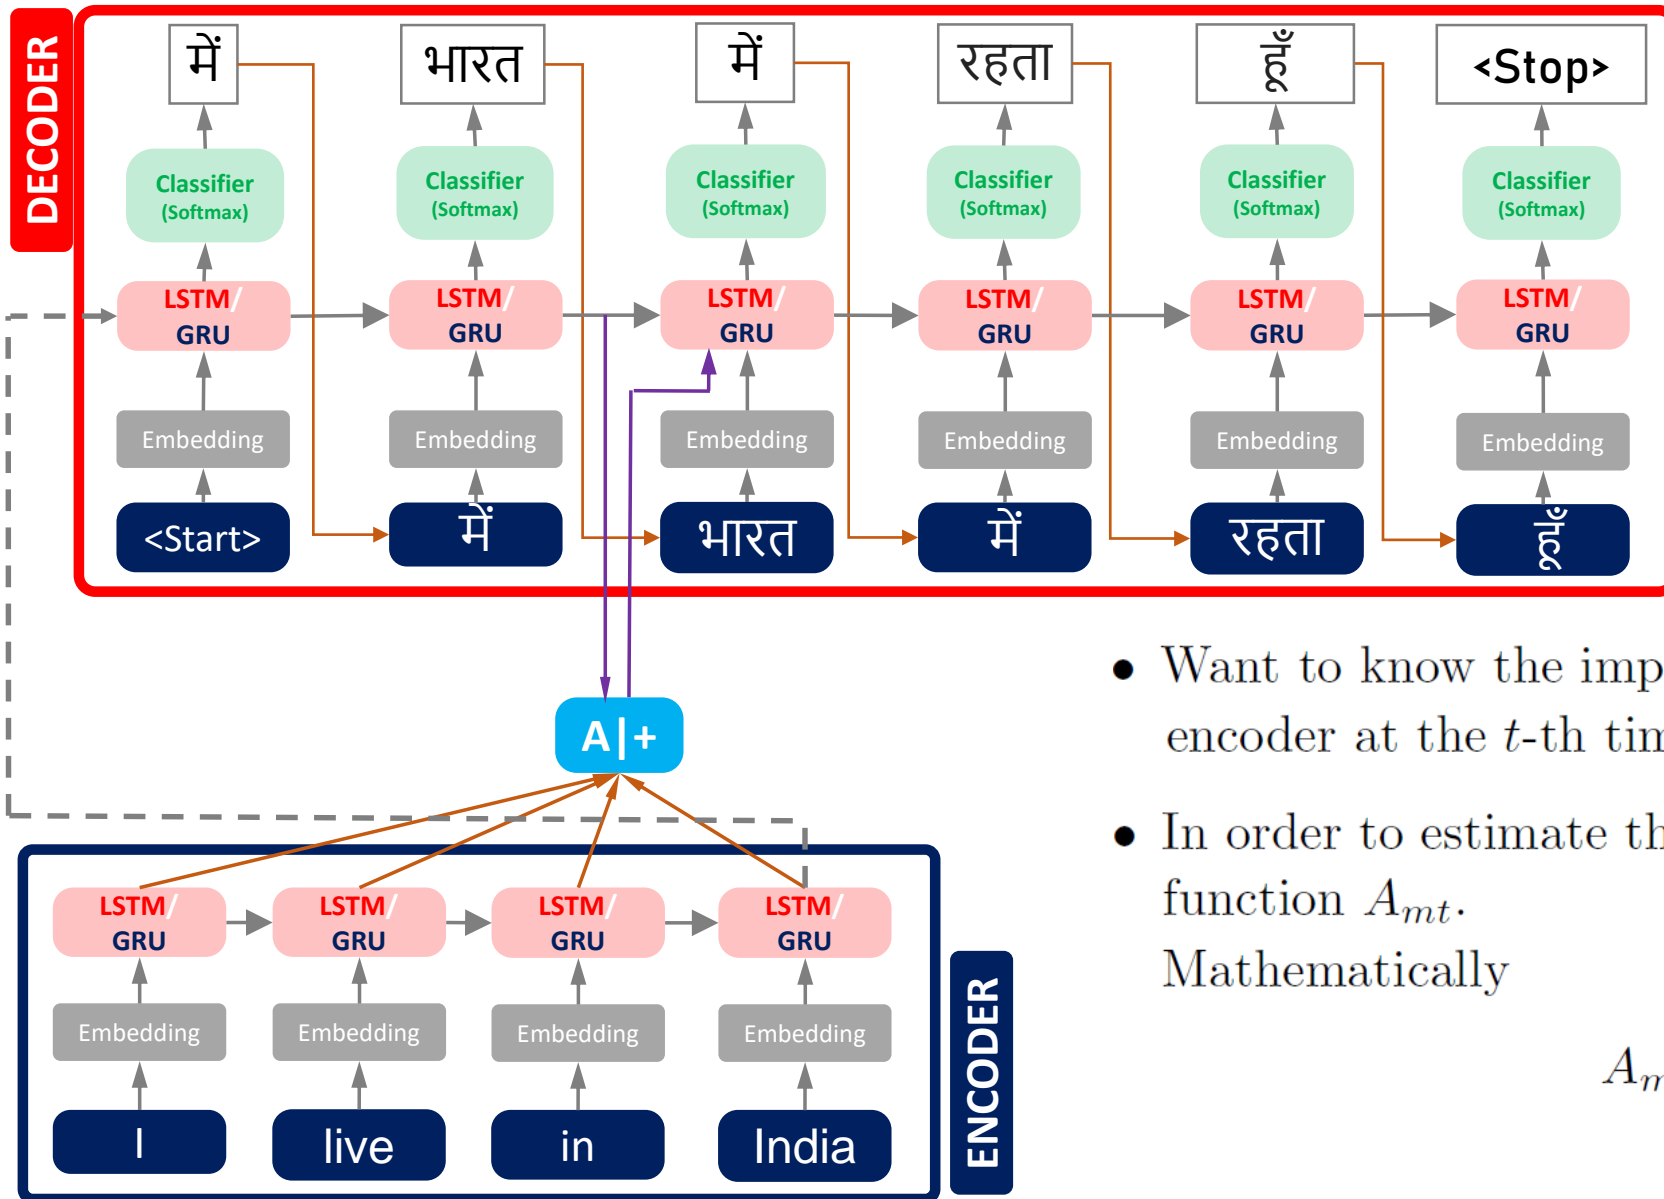


# Translation with Attention



- Want to know the importance of the  $m$ -th word of the encoder at the  $t$ -th time-step of the decoder:  $\beta_{mt}$ .
- In order to estimate the importance, we define a score function  $A_{mt}$ .  
Intuitively this function can be taken to depend upon
  - $\mathbf{h}_m^E$ : the encoder representation of the  $m$ -th word
  - $\mathbf{h}_{t-1}^D$ : the current decoder state

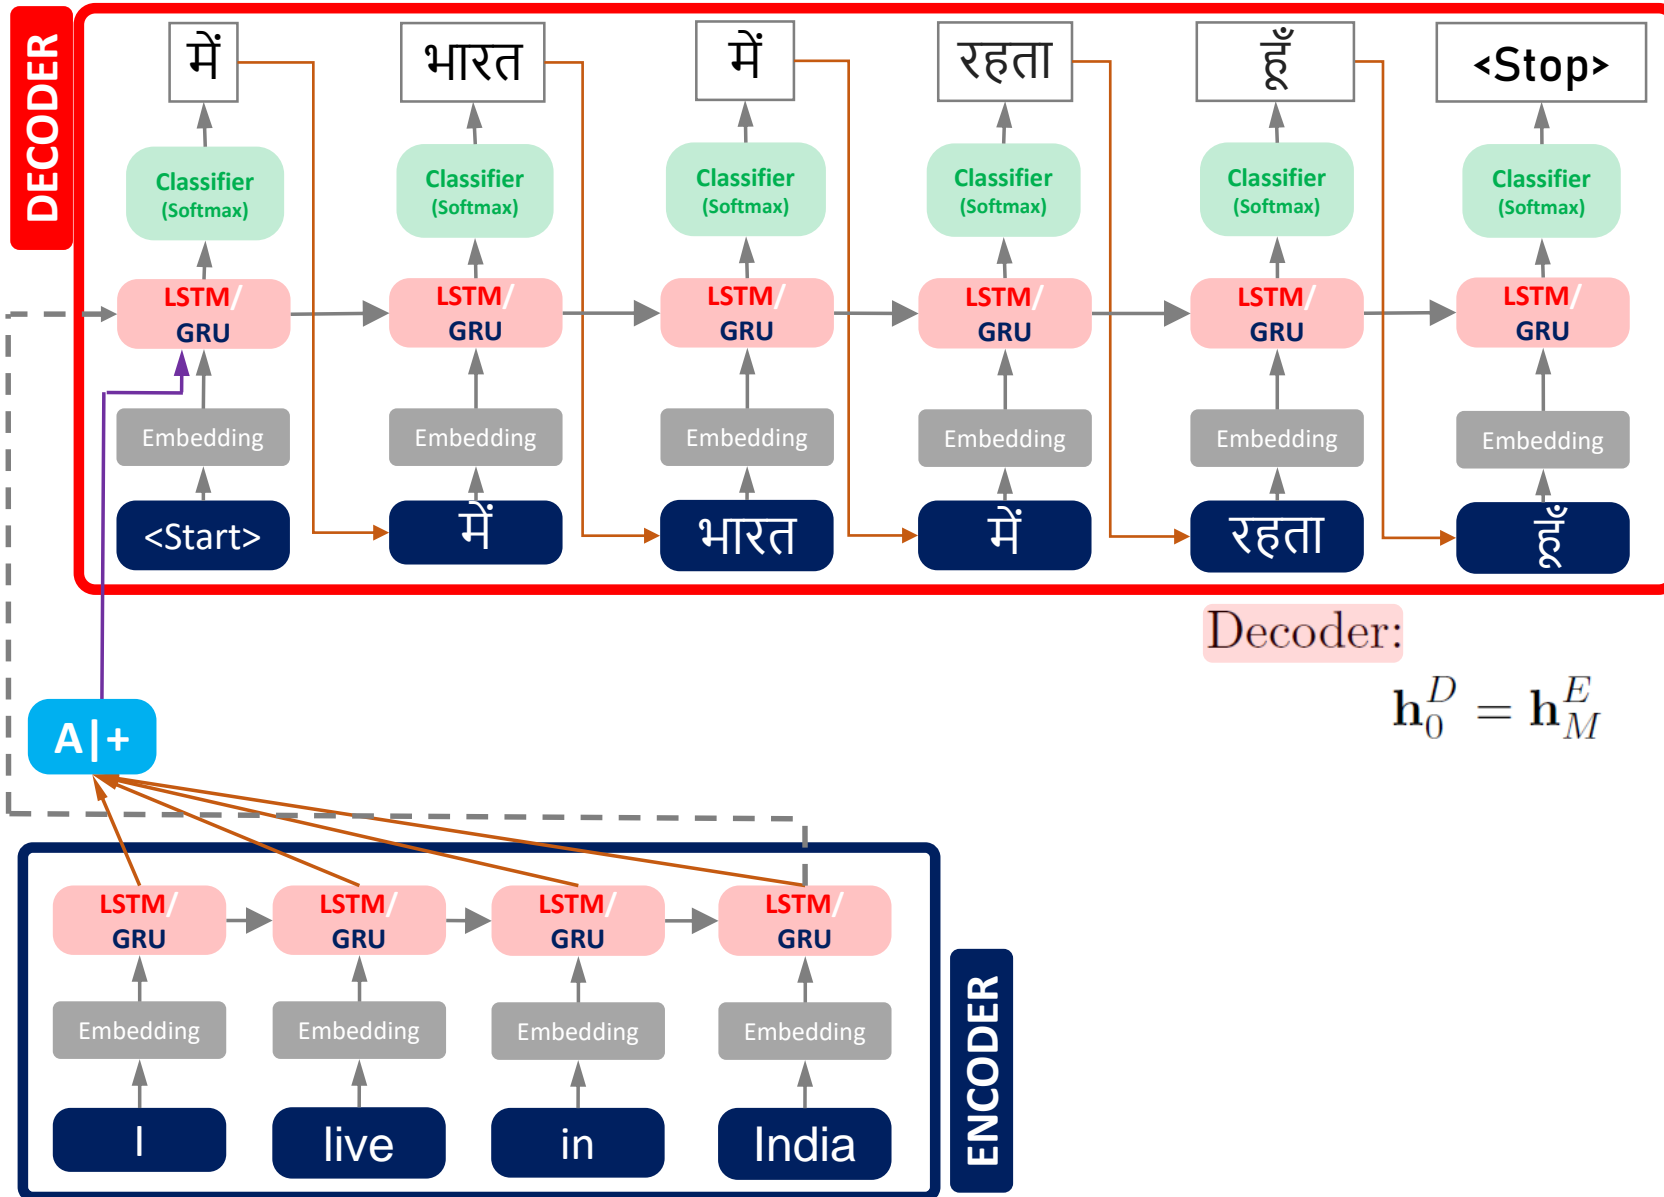
# Translation with Attention



- Want to know the importance of the  $m$ -th word of the encoder at the  $t$ -th time-step of the decoder:  $\beta_{mt}$ .
- In order to estimate the importance, we define a score function  $A_{mt}$ .  
Mathematically

$$A_{mt} = f(\mathbf{h}_m^E, \mathbf{h}_{t-1}^D)$$

# Translation with Attention



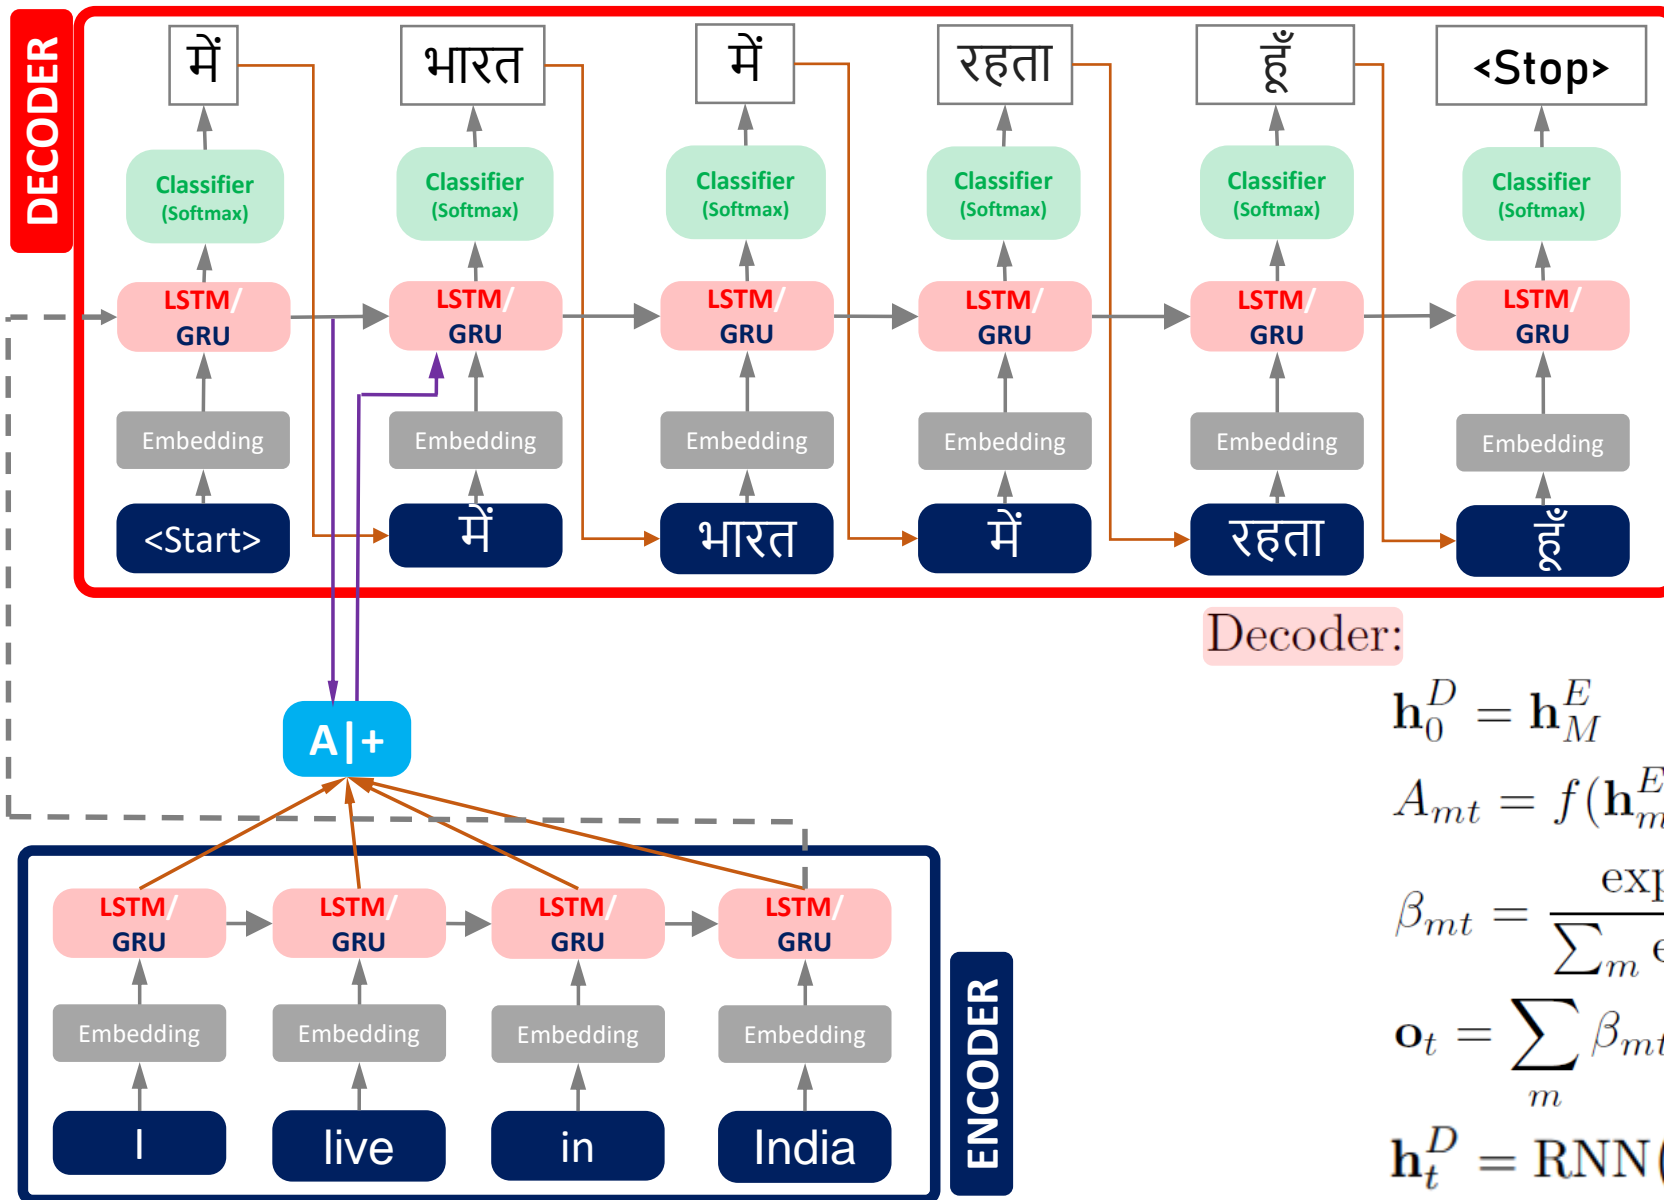
Encoder:

$$\mathbf{h}_m^E = \text{RNN}(\text{Embed}(\mathbf{x}_m), \mathbf{h}_{m-1}^E)$$

Decoder:

$$\mathbf{h}_0^D = \mathbf{h}_M^E$$

# Translation with Attention



Encoder:

$$\mathbf{h}_m^E = \text{RNN}(\text{Embed}(\mathbf{x}_m), \mathbf{h}_{m-1}^E)$$

Decoder:

$$\mathbf{h}_0^D = \mathbf{h}_M^E$$

$$A_{mt} = f(\mathbf{h}_m^E, \mathbf{h}_{t-1}^D)$$

Score function

$$\beta_{mt} = \frac{\exp(A_{mt})}{\sum_m \exp(A_{mt})}$$

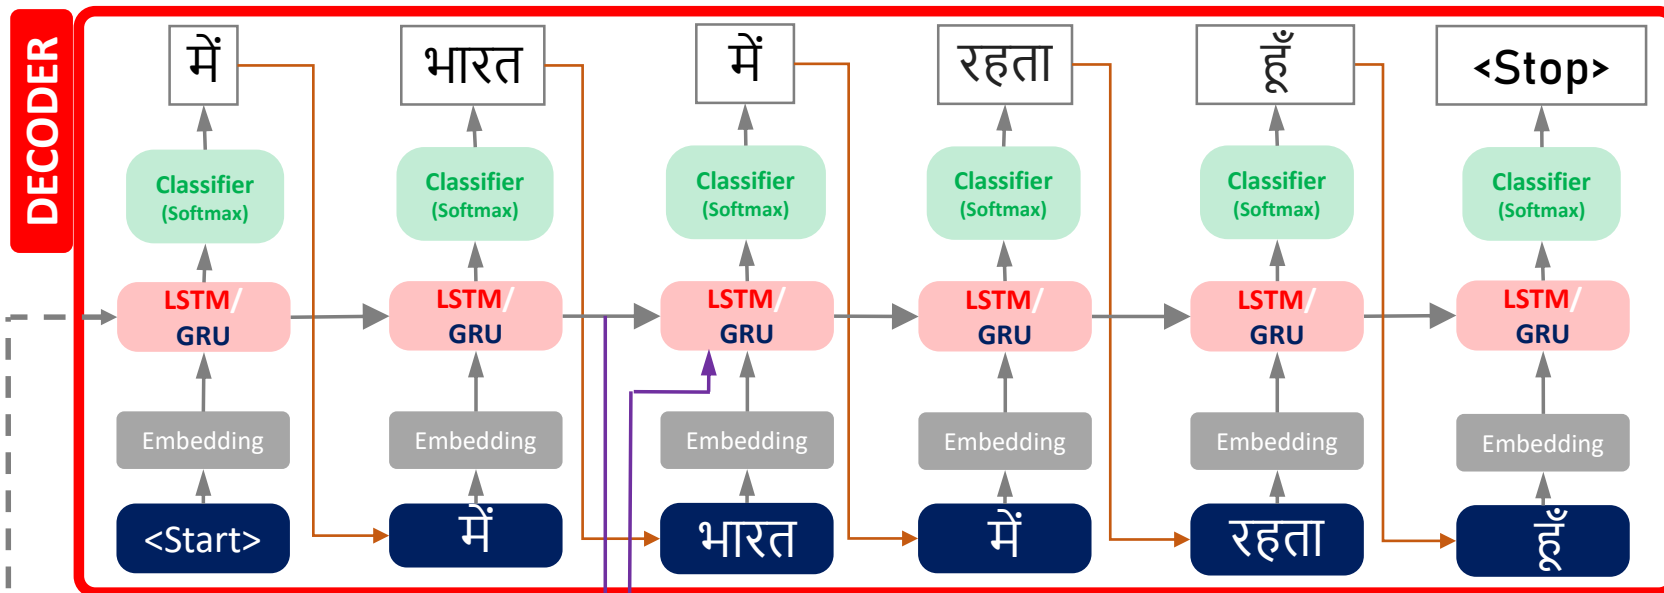
Attention weights

$$\mathbf{o}_t = \sum_m \beta_{mt} \mathbf{h}_m^E$$

Attention layer outputs

$$\mathbf{h}_t^D = \text{RNN}([\mathbf{o}_t, \text{Embed}(y_{t-1}^*)], \mathbf{h}_{t-1}^D)$$

# Translation with Attention



Encoder:

$$\mathbf{h}_m^E = \text{RNN}(\text{Embed}(\mathbf{x}_m), \mathbf{h}_{m-1}^E)$$

Decoder:

$$\mathbf{h}_0^D = \mathbf{h}_M^E$$

$$A_{mt} = f(\mathbf{h}_m^E, \mathbf{h}_{t-1}^D)$$

Score function

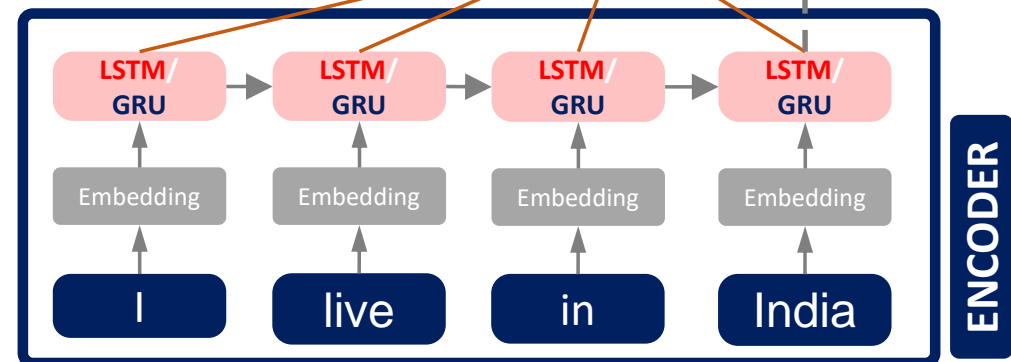
$$\beta_{mt} = \frac{\exp(A_{mt})}{\sum_m \exp(A_{mt})}$$

Attention weights

$$\mathbf{o}_t = \sum_m \beta_{mt} \mathbf{h}_m^E$$

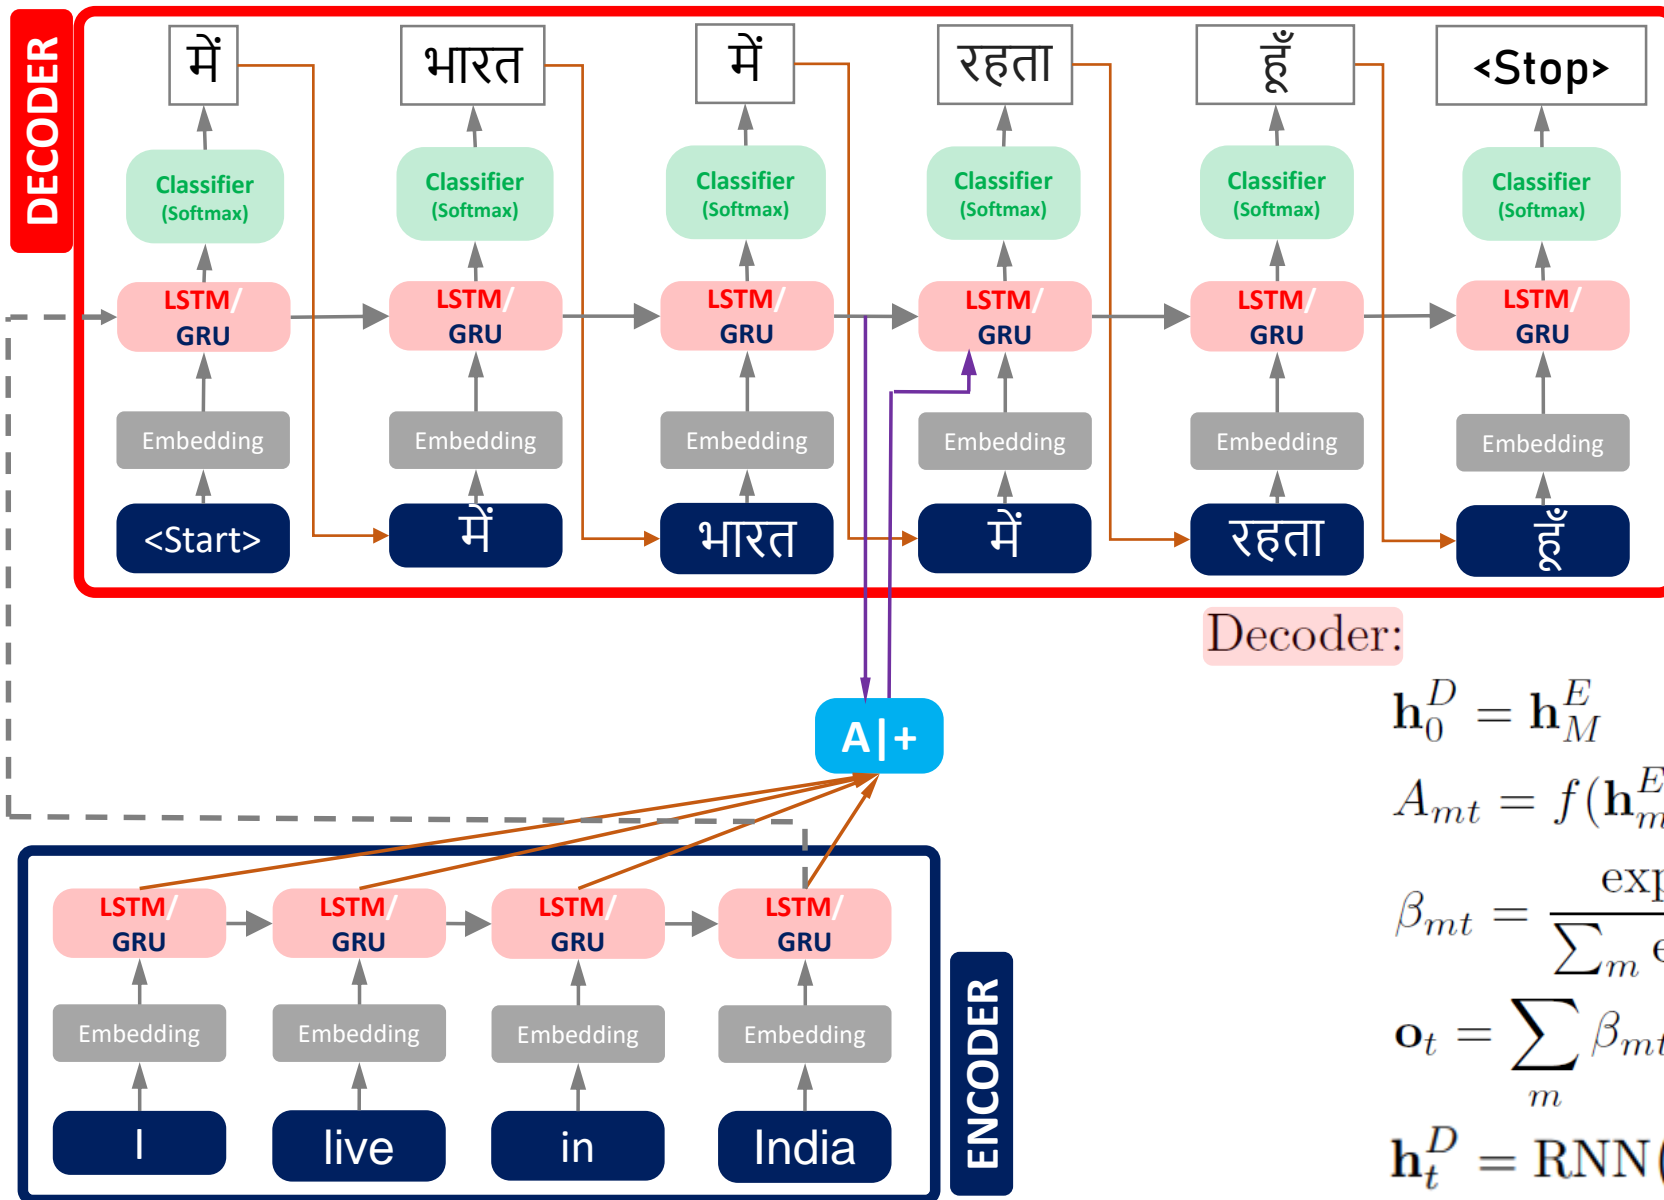
Attention layer outputs

$$\mathbf{h}_t^D = \text{RNN}([\mathbf{o}_t, \text{Embed}(y_{t-1}^*)], \mathbf{h}_{t-1}^D)$$





# Translation with Attention



Encoder:

$$\mathbf{h}_m^E = \text{RNN}(\text{Embed}(\mathbf{x}_m), \mathbf{h}_{m-1}^E)$$

Decoder:

$$\mathbf{h}_0^D = \mathbf{h}_M^E$$

$$A_{mt} = f(\mathbf{h}_m^E, \mathbf{h}_{t-1}^D)$$

Score function

$$\beta_{mt} = \frac{\exp(A_{mt})}{\sum_m \exp(A_{mt})}$$

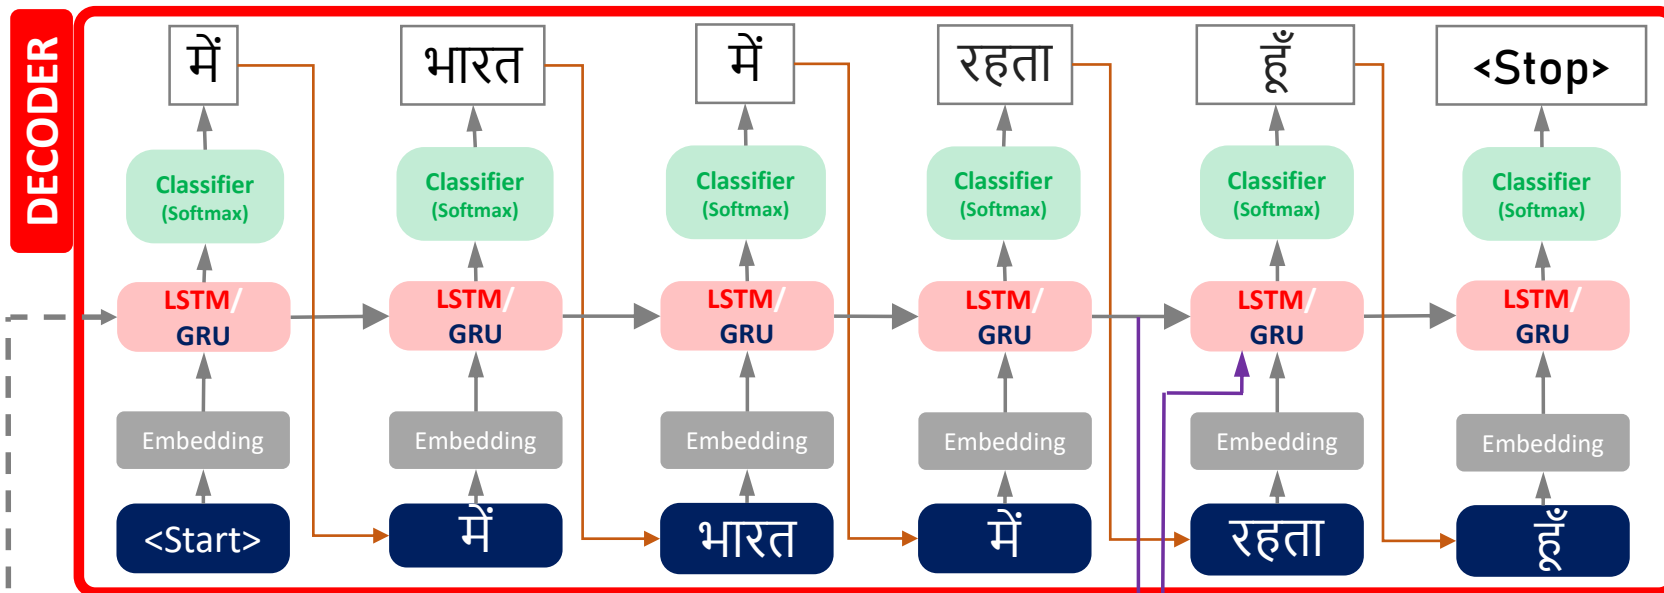
Attention weights

$$\mathbf{o}_t = \sum_m \beta_{mt} \mathbf{h}_m^E$$

Attention layer outputs

$$\mathbf{h}_t^D = \text{RNN}([\mathbf{o}_t, \text{Embed}(y_{t-1}^*)], \mathbf{h}_{t-1}^D)$$

# Translation with Attention



Encoder:

$$\mathbf{h}_m^E = \text{RNN}(\text{Embed}(\mathbf{x}_m), \mathbf{h}_{m-1}^E)$$

Decoder:

$$\mathbf{h}_0^D = \mathbf{h}_M^E$$

$$A_{mt} = f(\mathbf{h}_m^E, \mathbf{h}_{t-1}^D)$$

Score function

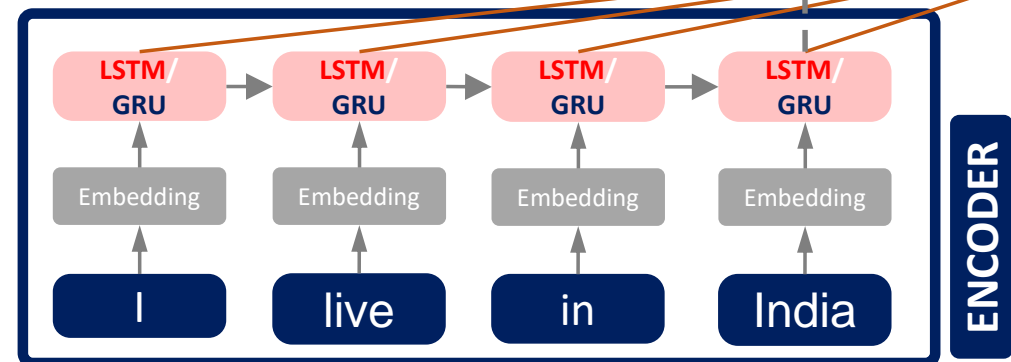
$$\beta_{mt} = \frac{\exp(A_{mt})}{\sum_m \exp(A_{mt})}$$

Attention weights

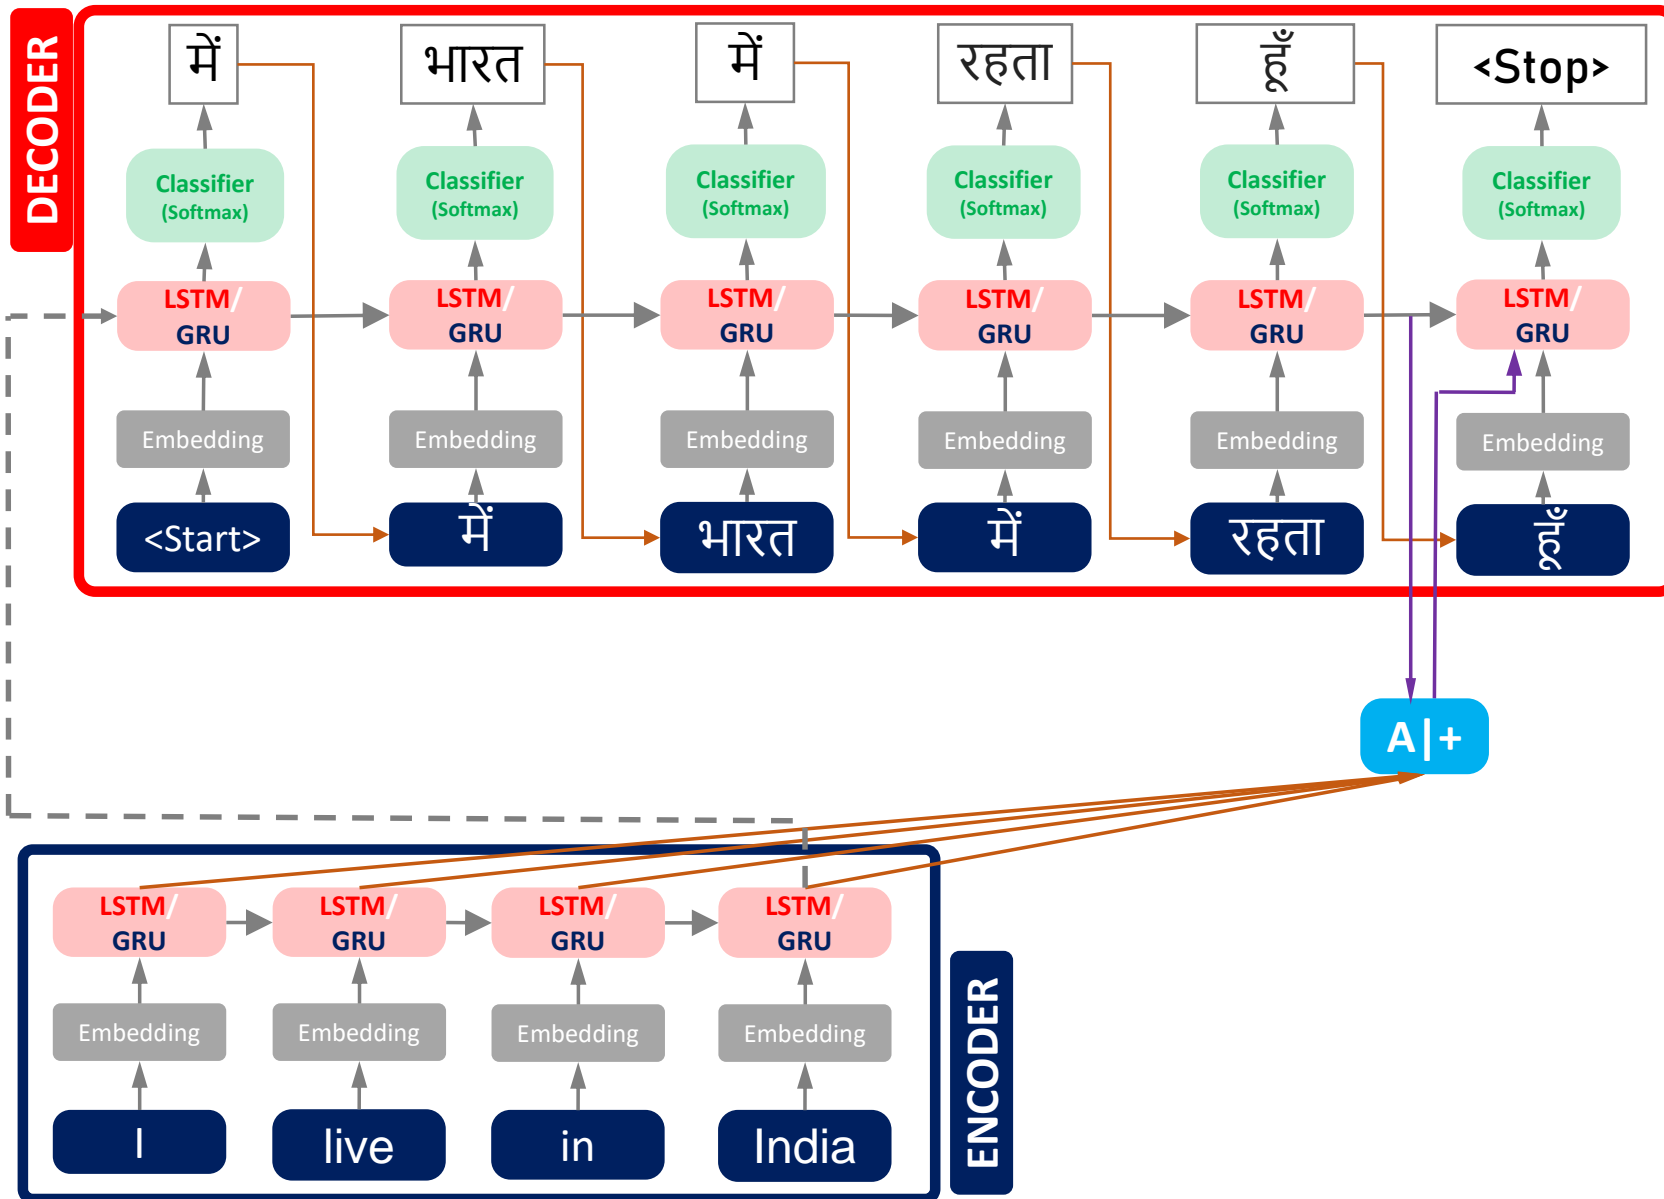
$$\mathbf{o}_t = \sum_m \beta_{mt} \mathbf{h}_m^E$$

Attention layer outputs

$$\mathbf{h}_t^D = \text{RNN}([\mathbf{o}_t, \text{Embed}(y_{t-1}^*)], \mathbf{h}_{t-1}^D)$$



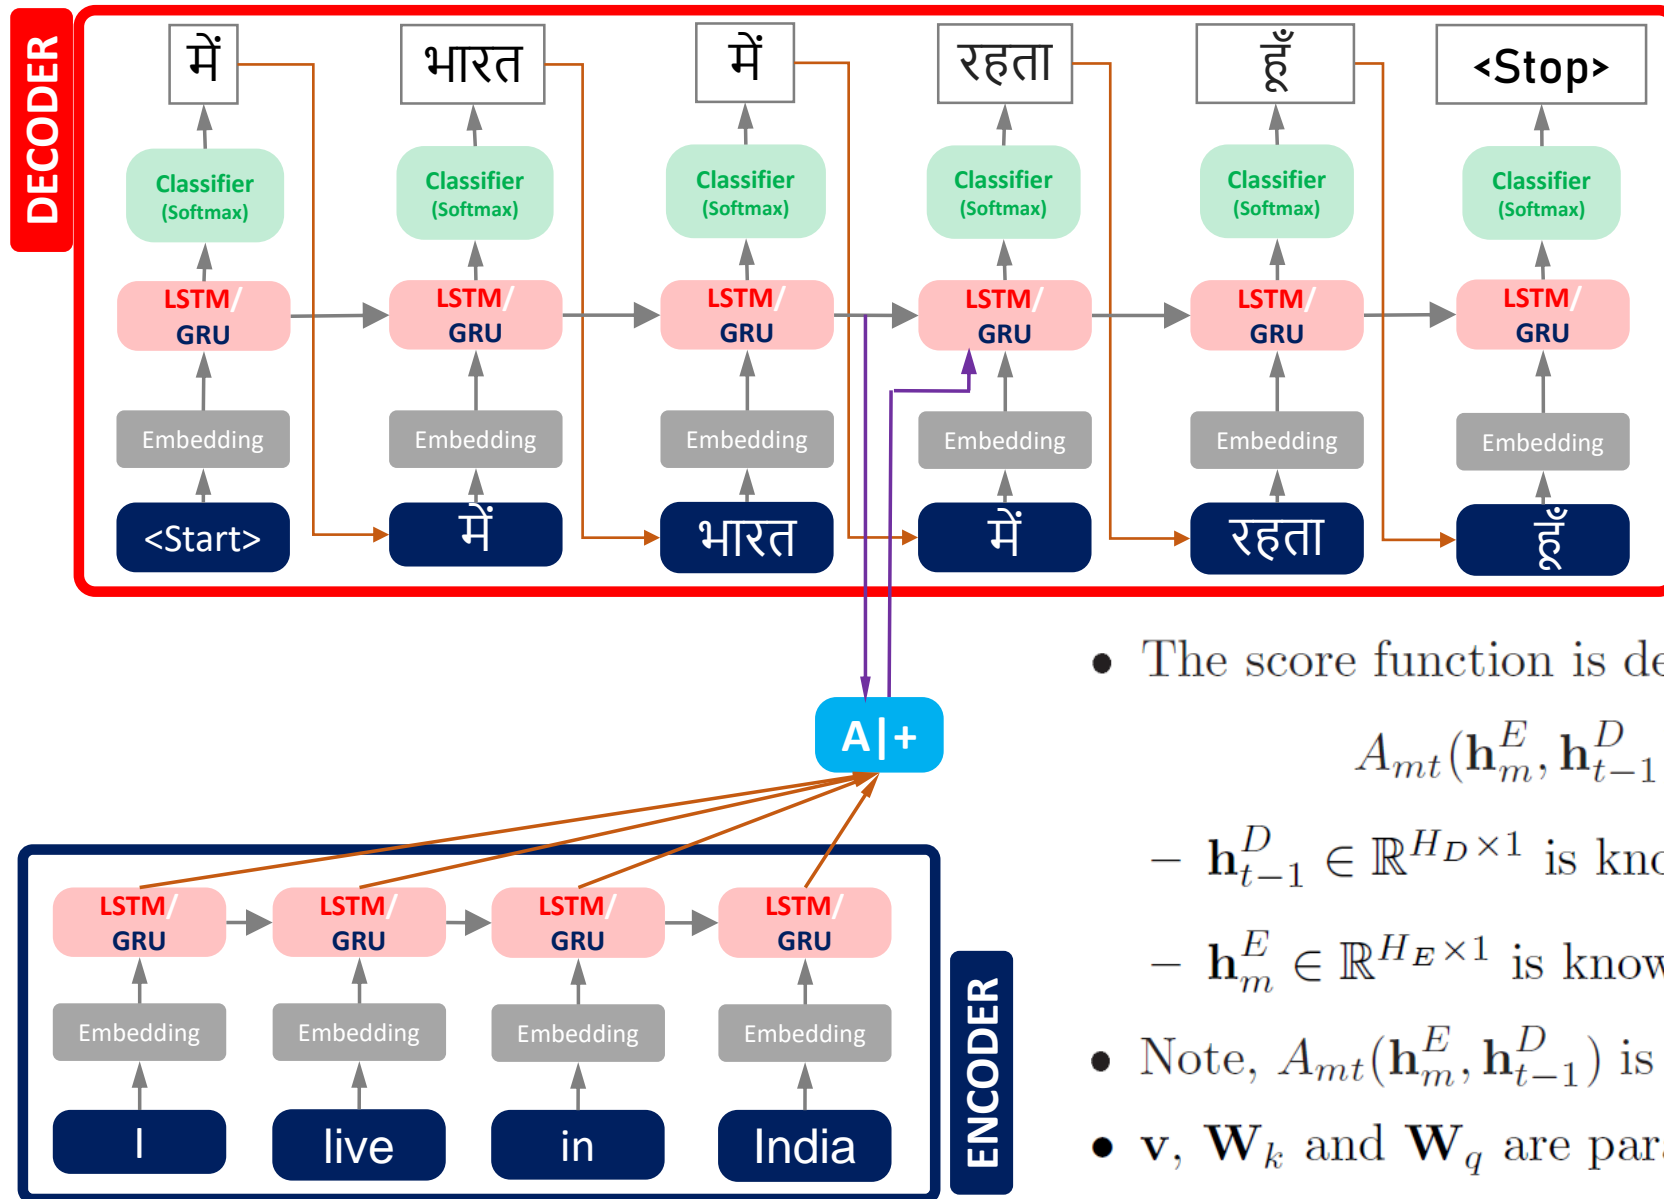
# Translation with Attention



# Masked softmax

- Often input sequences are padded.
- Masked softmax filters out those elements.
  - Need to specify the valid length of each sequence.
- Masked softmax assigns 0 weight to elements outside the valid length.
  - It takes those elements to be large negative, such that after softmax the weights become 0.

# Bahdanau's attention



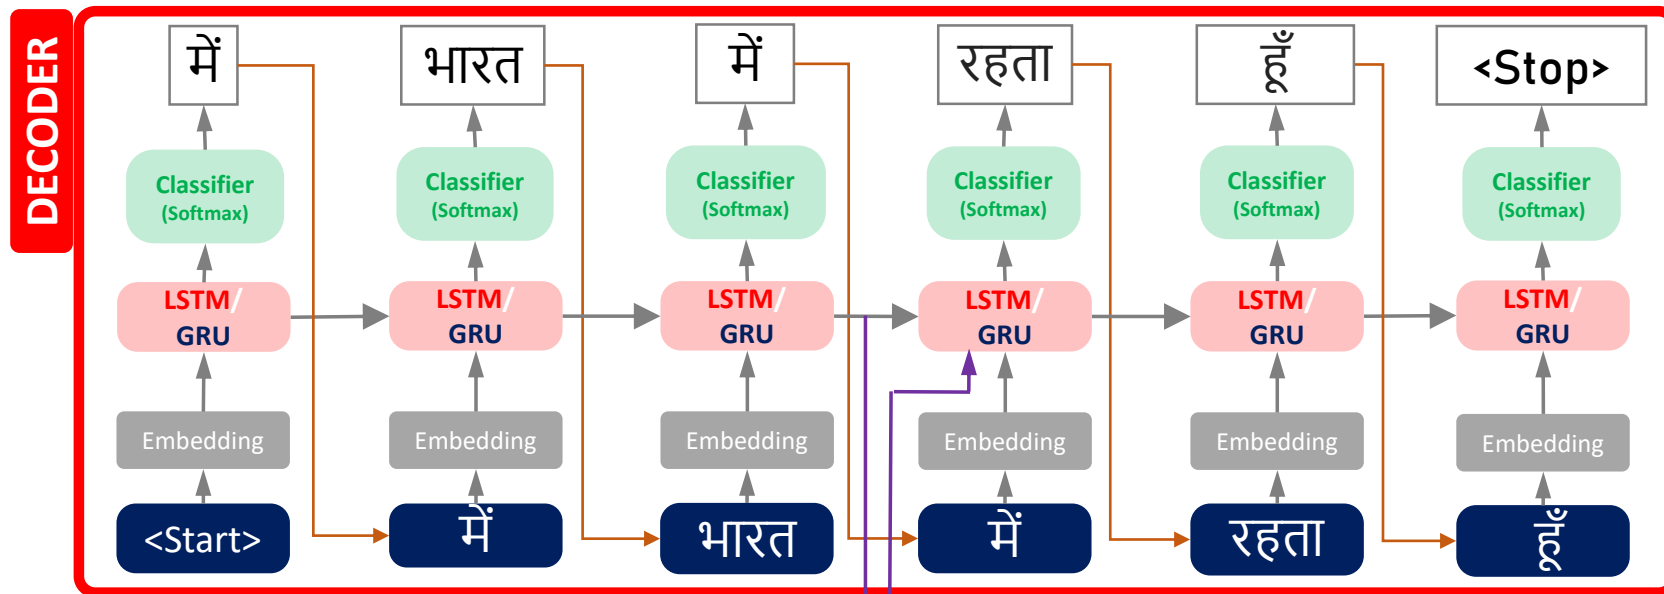
- The score function is defined as

$$A_{mt}(\mathbf{h}_m^E, \mathbf{h}_{t-1}^D) = \mathbf{v}^T \tanh(\mathbf{W}_k \mathbf{h}_m^E + \mathbf{W}_q \mathbf{h}_{t-1}^D)$$

- $\mathbf{h}_{t-1}^D \in \mathbb{R}^{H_D \times 1}$  is known as the query  $\mathbf{q}$ .
- $\mathbf{h}_m^E \in \mathbb{R}^{H_E \times 1}$  is known as the key  $\mathbf{k}$ .
- Note,  $A_{mt}(\mathbf{h}_m^E, \mathbf{h}_{t-1}^D)$  is scalar.
- $\mathbf{v}$ ,  $\mathbf{W}_k$  and  $\mathbf{W}_q$  are parameters that are learnt.



# Bahdanau's attention



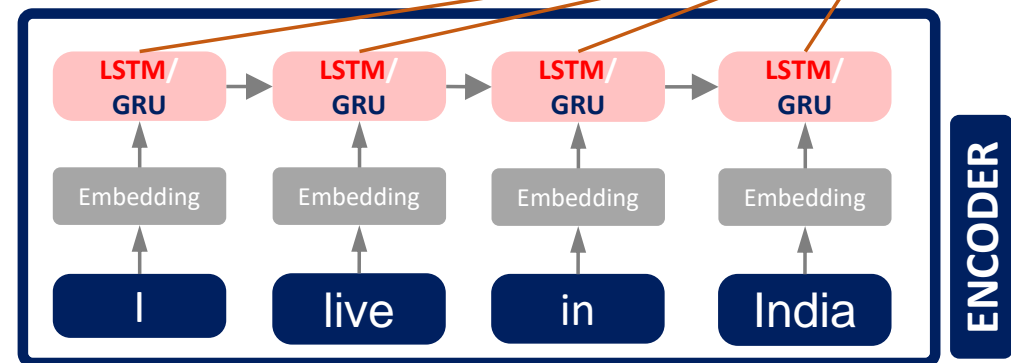
- The softmax function is then used to obtain the attention weights

$$\beta_{mt} = \frac{\exp(A_{mt})}{\sum_m \exp(A_{mt})}$$

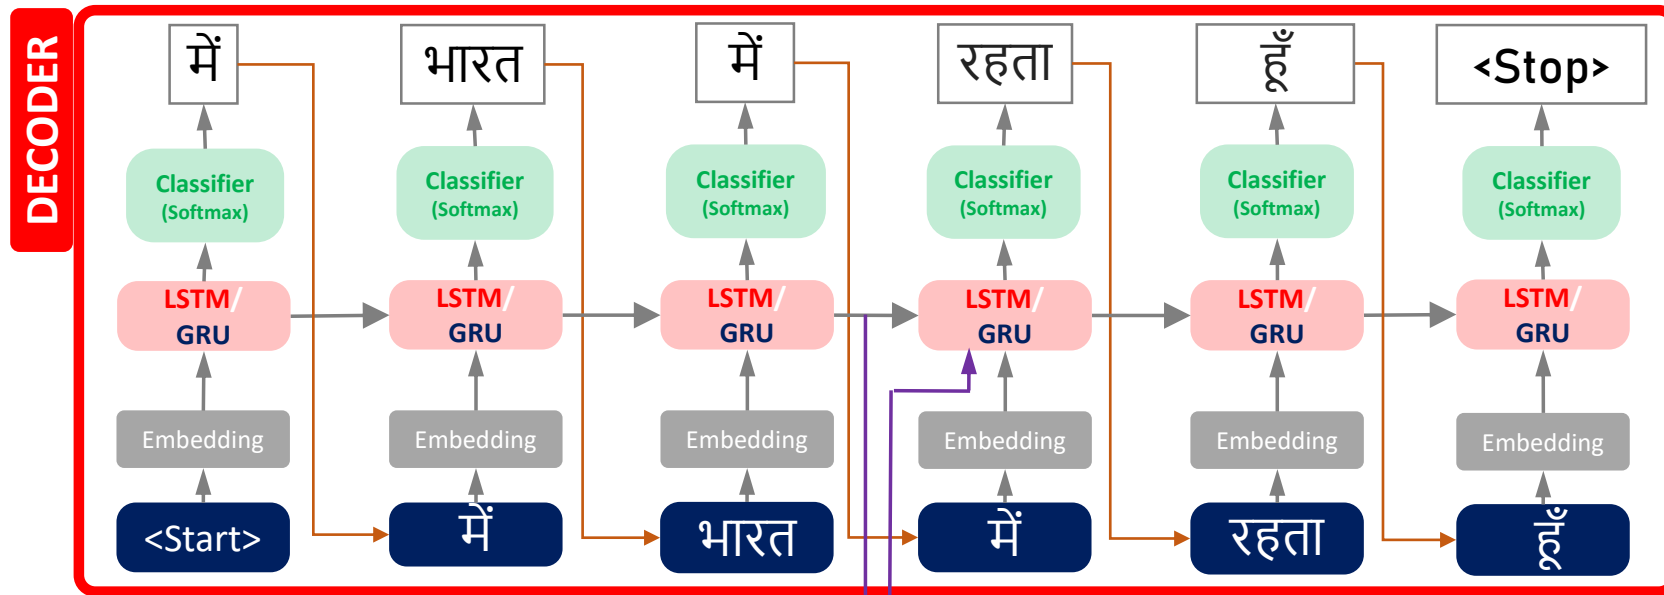
where  $\beta_{mt}$  is the weight given to the  $m$ th input word at the  $t$ th time-step of the decoder.

- The output of attention layer is the weighted sum of the values

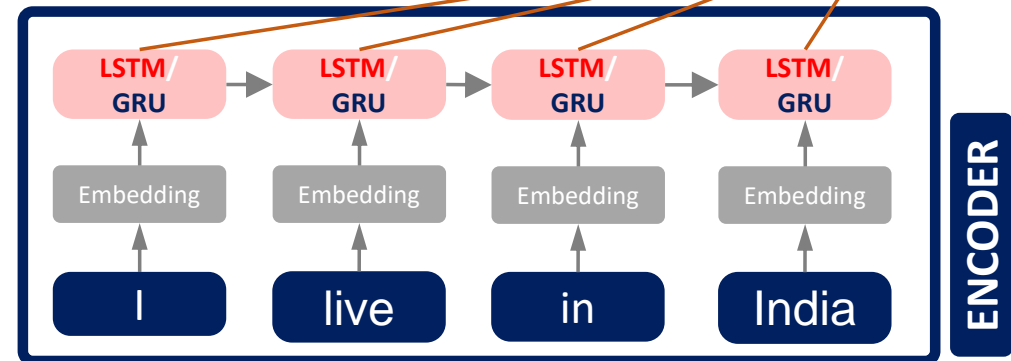
$$\mathbf{o}_t = \sum_m \beta_{mt} h_m^E$$



# Dot-product attention



A|+



- The score function is defined as the dot-product of the query and a key, and divided by the square-root of the dimension:

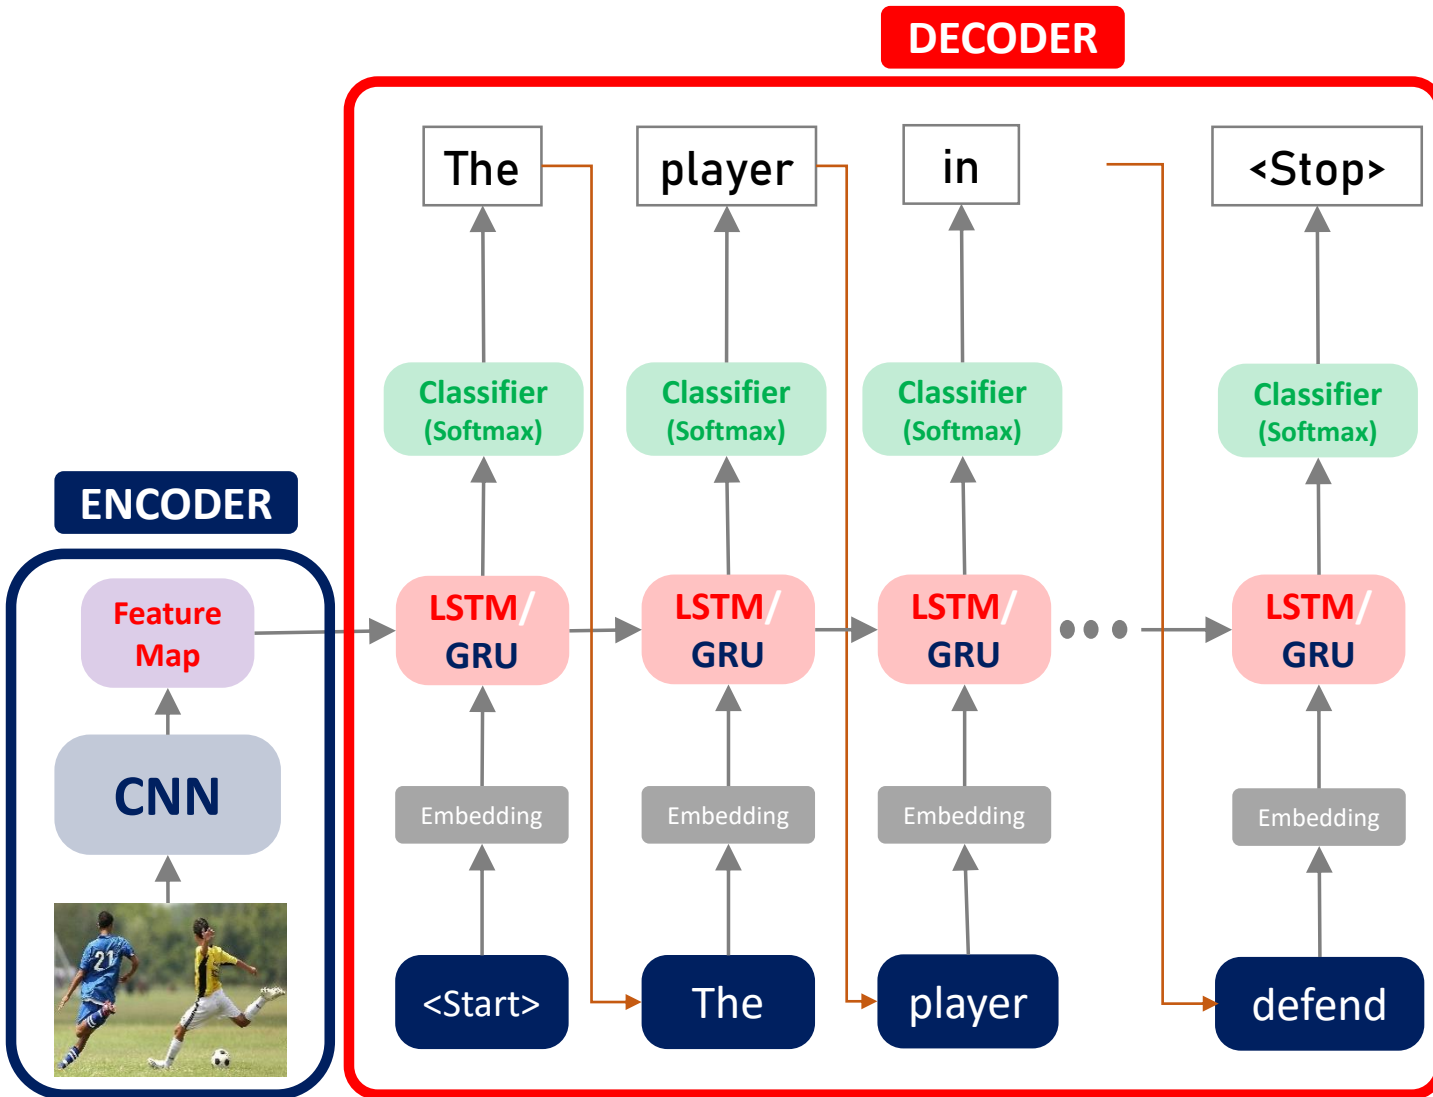
$$A(\mathbf{q}, \mathbf{k}) = \frac{\langle \mathbf{q}, \mathbf{k} \rangle}{\sqrt{d}}$$

– Here  $d$  is the dimension of the key vector.

- The scaling of the dot products by  $1/\sqrt{d}$  is done to facilitate achieving stable gradients.


# VISUAL ATTENTION

# Image captioning: standard model



- A key aspect of the human visual system is attention.
- The decoder model presented here uses a static representation of the image.
- Attention mechanism enables salient features to play important roles when needed.


# CNN feature map



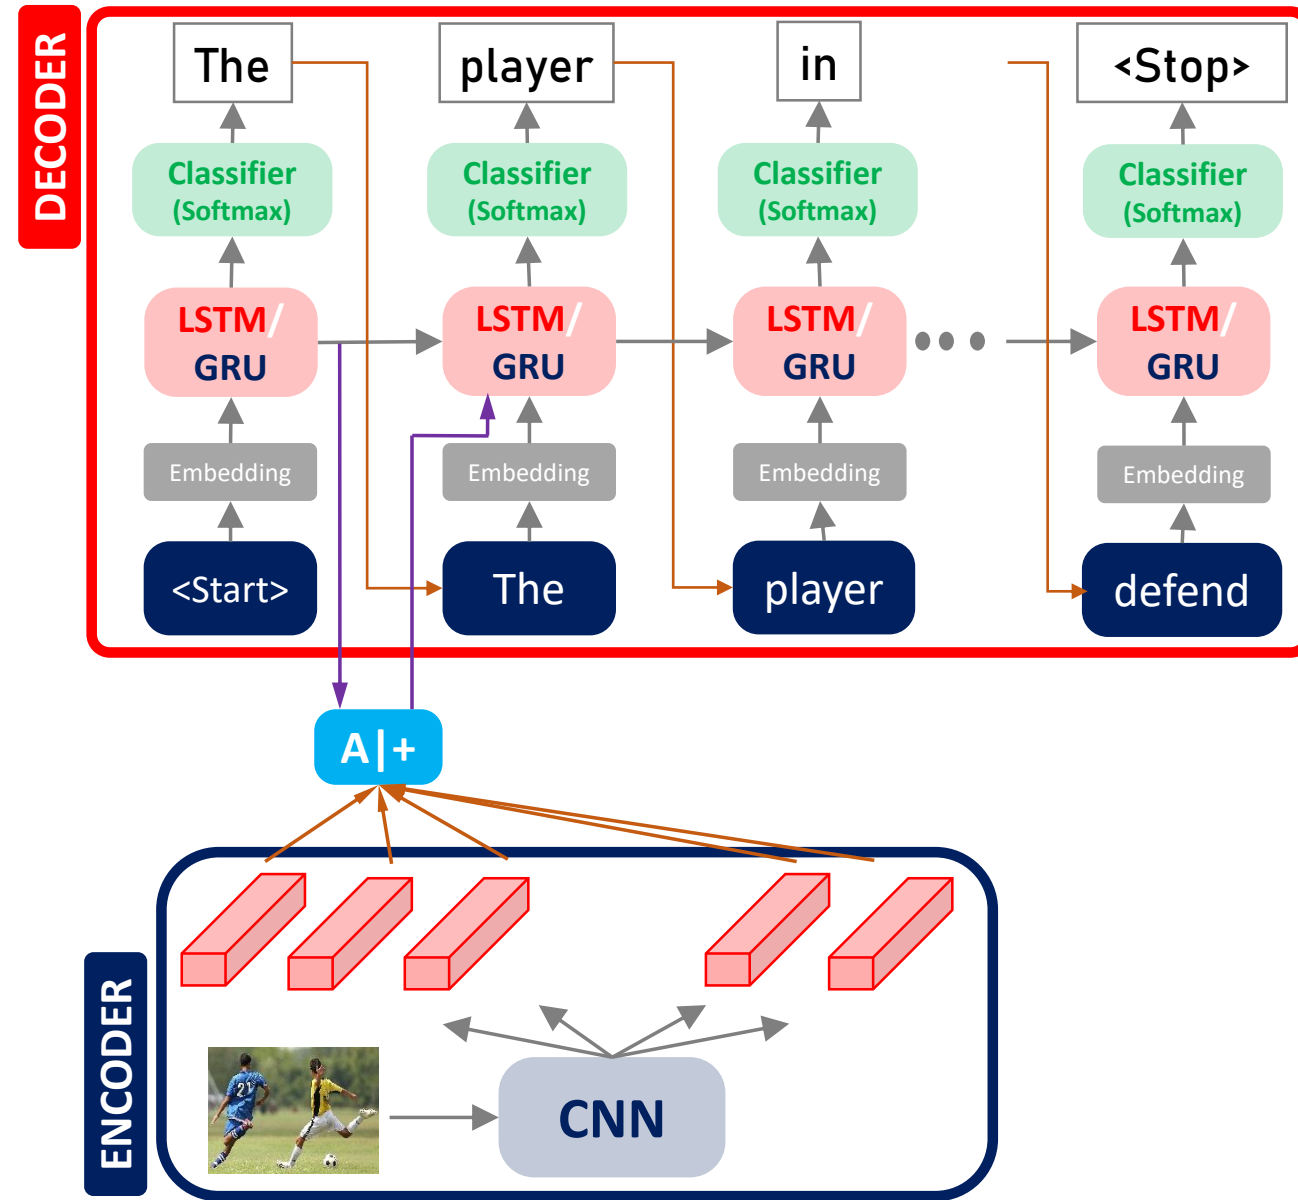
- The feature map from a FC layer represent information in a very compact form.
  - This can lead to loss of useful information.
  - Also spatial information is not properly retained.
- Output is taken from one of the convolutional layers.
  - This is the feature map generated in one of the convolutional layers.



# CNN feature map



# Visual attention model



- Suppose  $\mathbf{f}_1, \mathbf{f}_2, \dots, \mathbf{f}_M$  are the feature vectors derived from a convolutional layer.
- Let  $\mathbf{y}^{*(t)}$  be the output of the decoder at time-step  $t$ .
- The hidden state at time-step  $t$  of the RNN can be computed as

$$\mathbf{h}_t = \text{RNN}(\mathbf{h}_{t-1}, [\mathbf{y}^{*(t-1)}, \mathbf{o}_t])$$

where  $\mathbf{o}_t$  is the weighted sum of the CNN feature vectors

$$\mathbf{o}_t = \sum_{m=1}^M \beta_{mt} \mathbf{f}_m$$