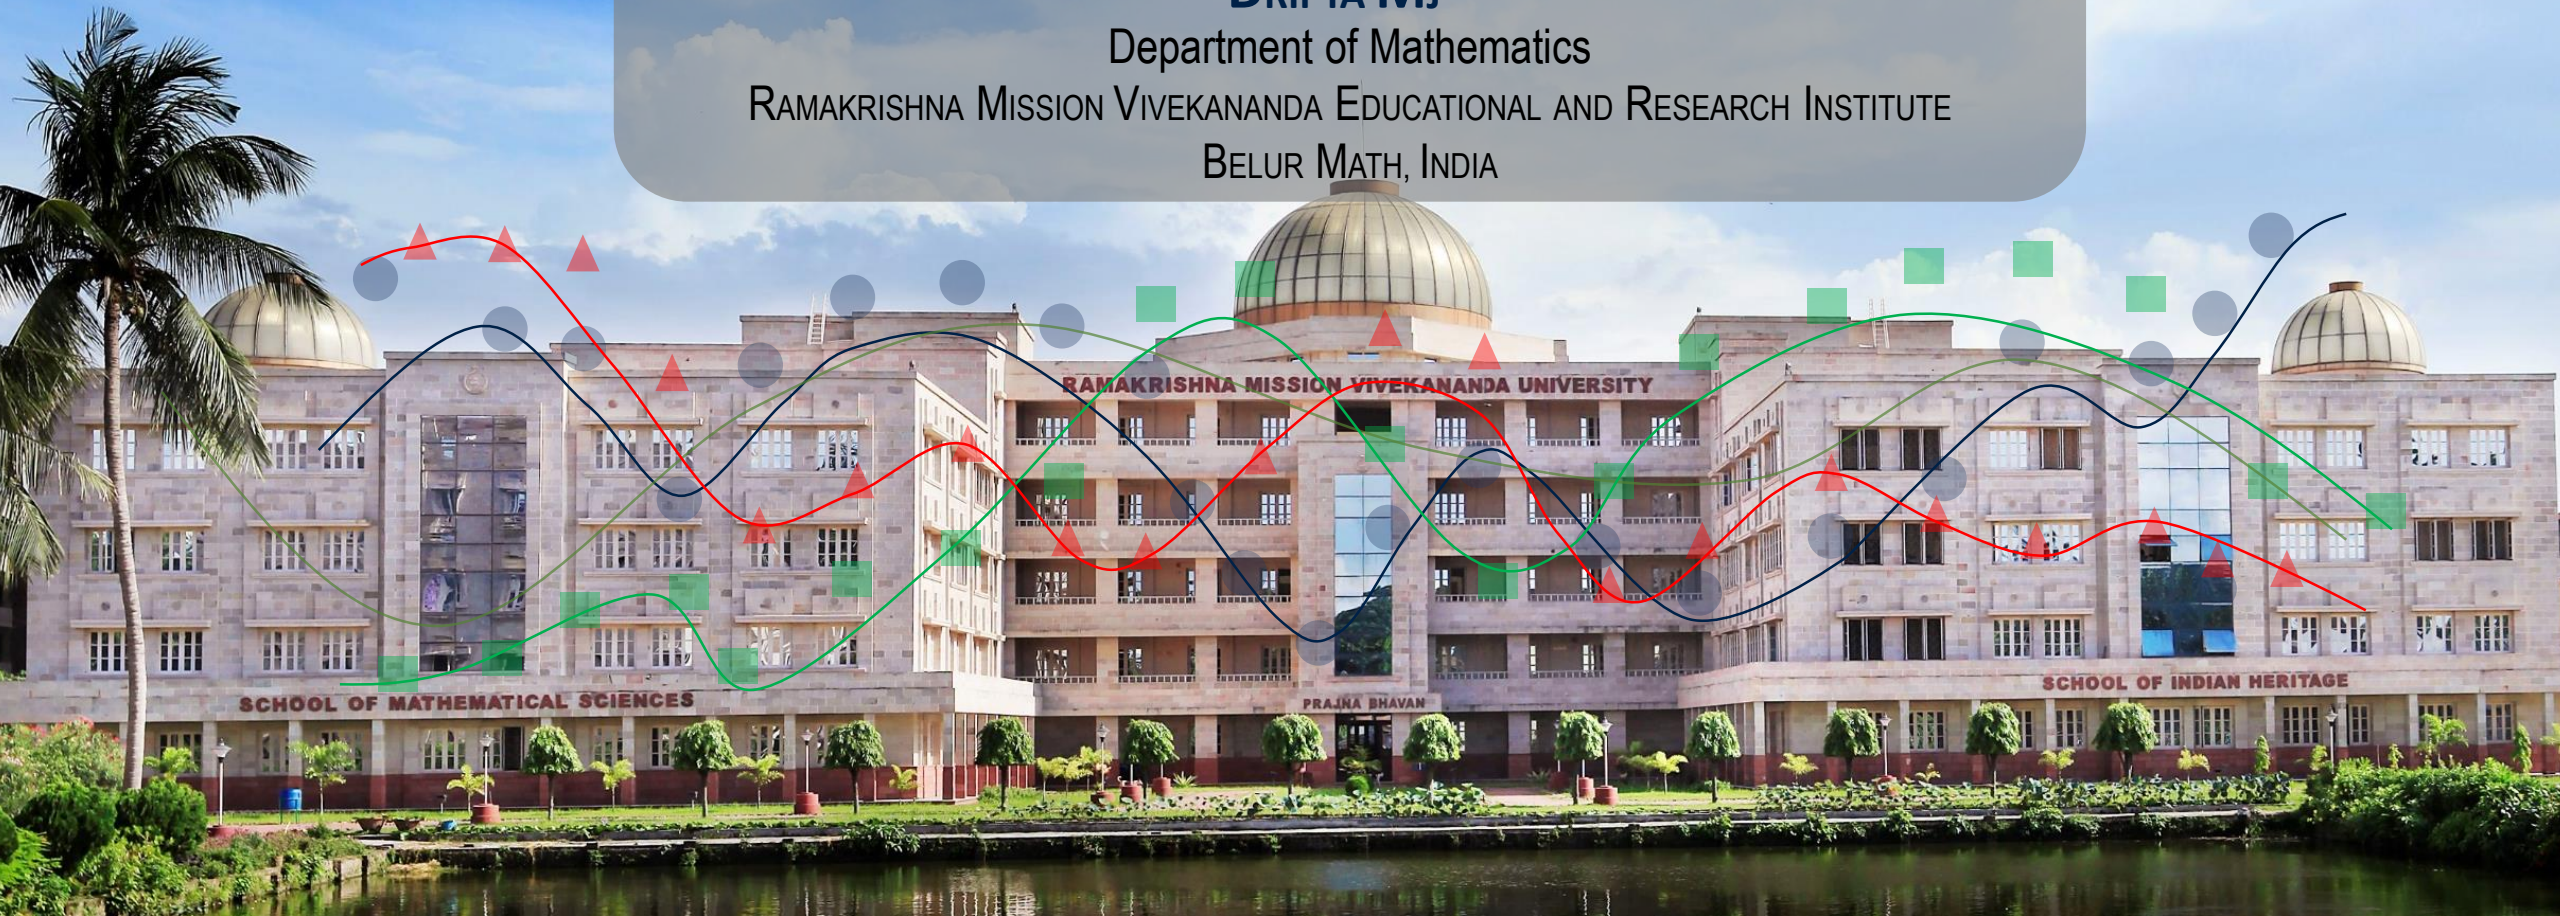


# Transformers

**DRIPTA MJ**

Department of Mathematics

RAMAKRISHNA MISSION VIVEKANANDA EDUCATIONAL AND RESEARCH INSTITUTE  
BELUR MATH, INDIA



# Transformer

- Issues with RNNs:
  - Cannot parallelize within a sequence due to recurrence.
  - Problem of vanishing/exploding gradients.
  - Requires large number of training steps.
- Vaswani et. al. 2017 proposed a new architecture based on the attention mechanism

## Attention Is All You Need

- The architecture is known as [Transformer](#).
- Employs a self-attention mechanism to model relationships between all words in a sentence



# Self-attention

- In typical encoder-decoder attention layers mechanisms, the queries come from the decoder layer, while the keys and values come from the output of the encoder.
  - This allows every position in the decoder to attend over all positions in the input sequence.
- Self-attention, sometimes called intra-attention is an attention mechanism relating different positions of a single sequence in order to compute a representation of the sequence.
- Suppose we have an encoder with self-attention layers.
  - In a self-attention layer, the keys, values and queries come from the output of the previous layer in the encoder.
- Self-attention layers in the decoder allow each position in the decoder to attend all positions in the decoder up to and including that position.

# Self-attention


- In each step, the transformer applies a self-attention mechanism which directly models relationships between all words in a sentence, regardless of their respective position.
- Consider the following sentence:

We need a crane to lift the item


- The model needs to determine that the word “crane” refers to the a large machine and not the bird.
- The Transformer computes the next representation of each word by comparing it with every other word in the sentence.

# Self-attention

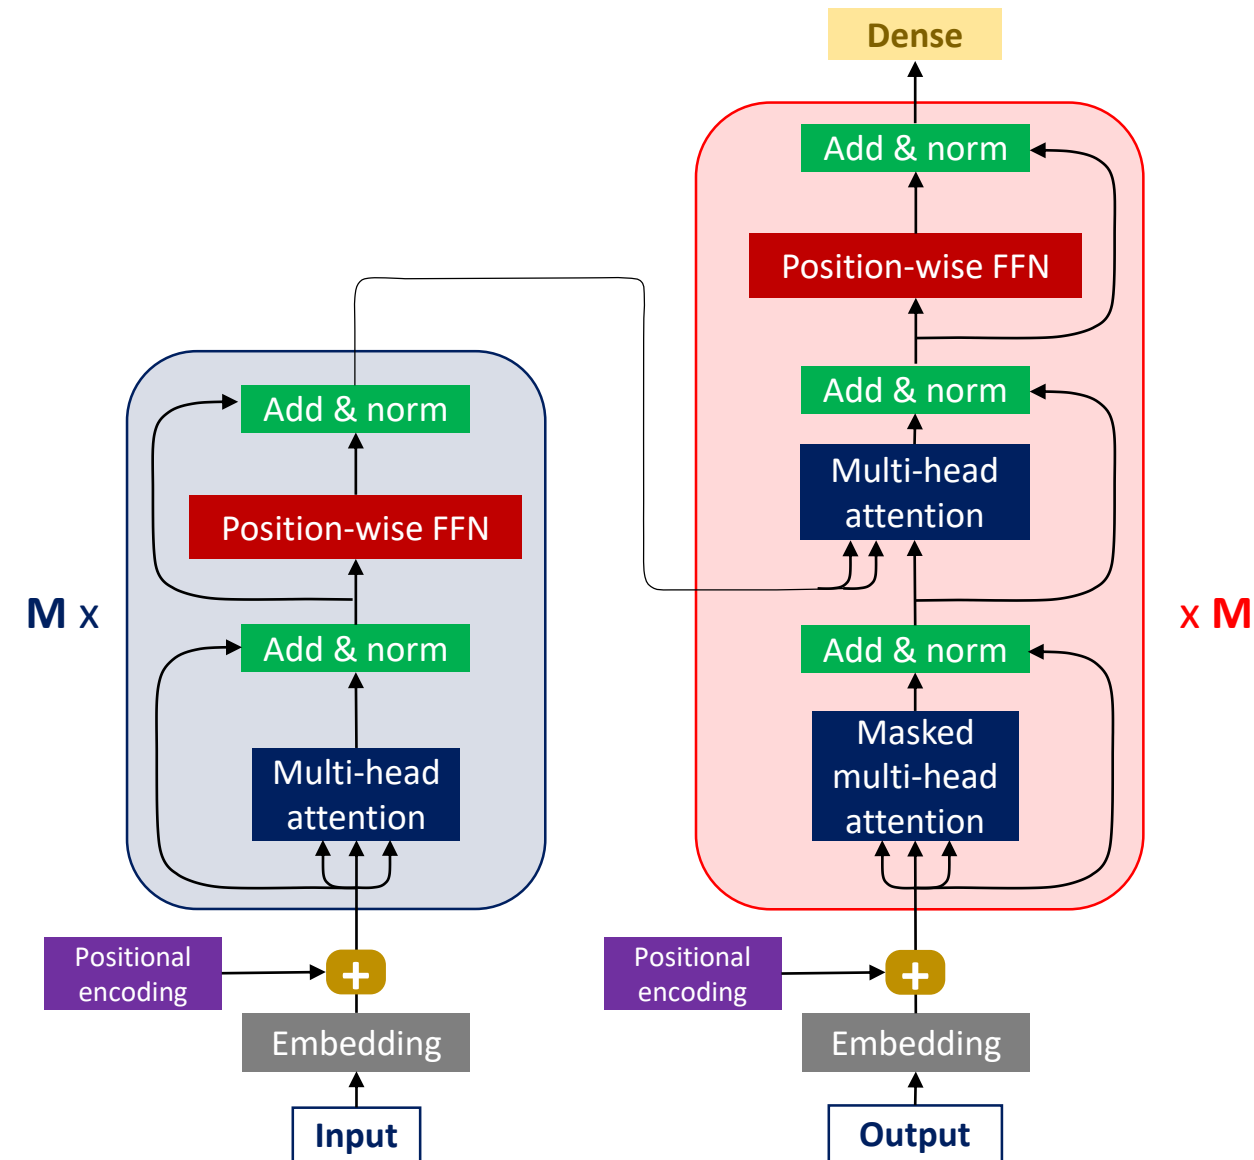
## ENCODER SELF-ATTENTION



## DECODER SELF-ATTENTION (MASKED)



# Transformer Architecture

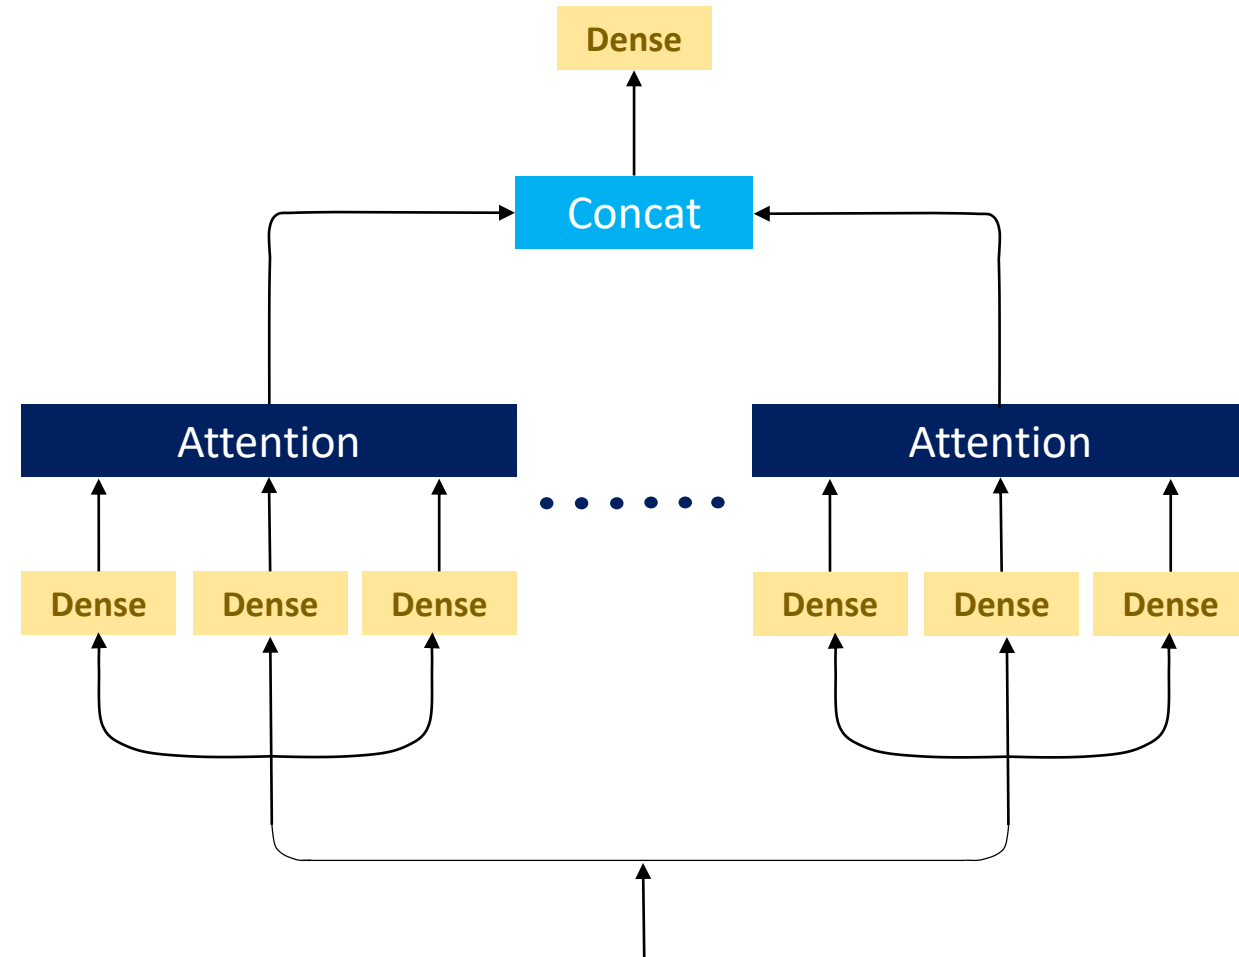
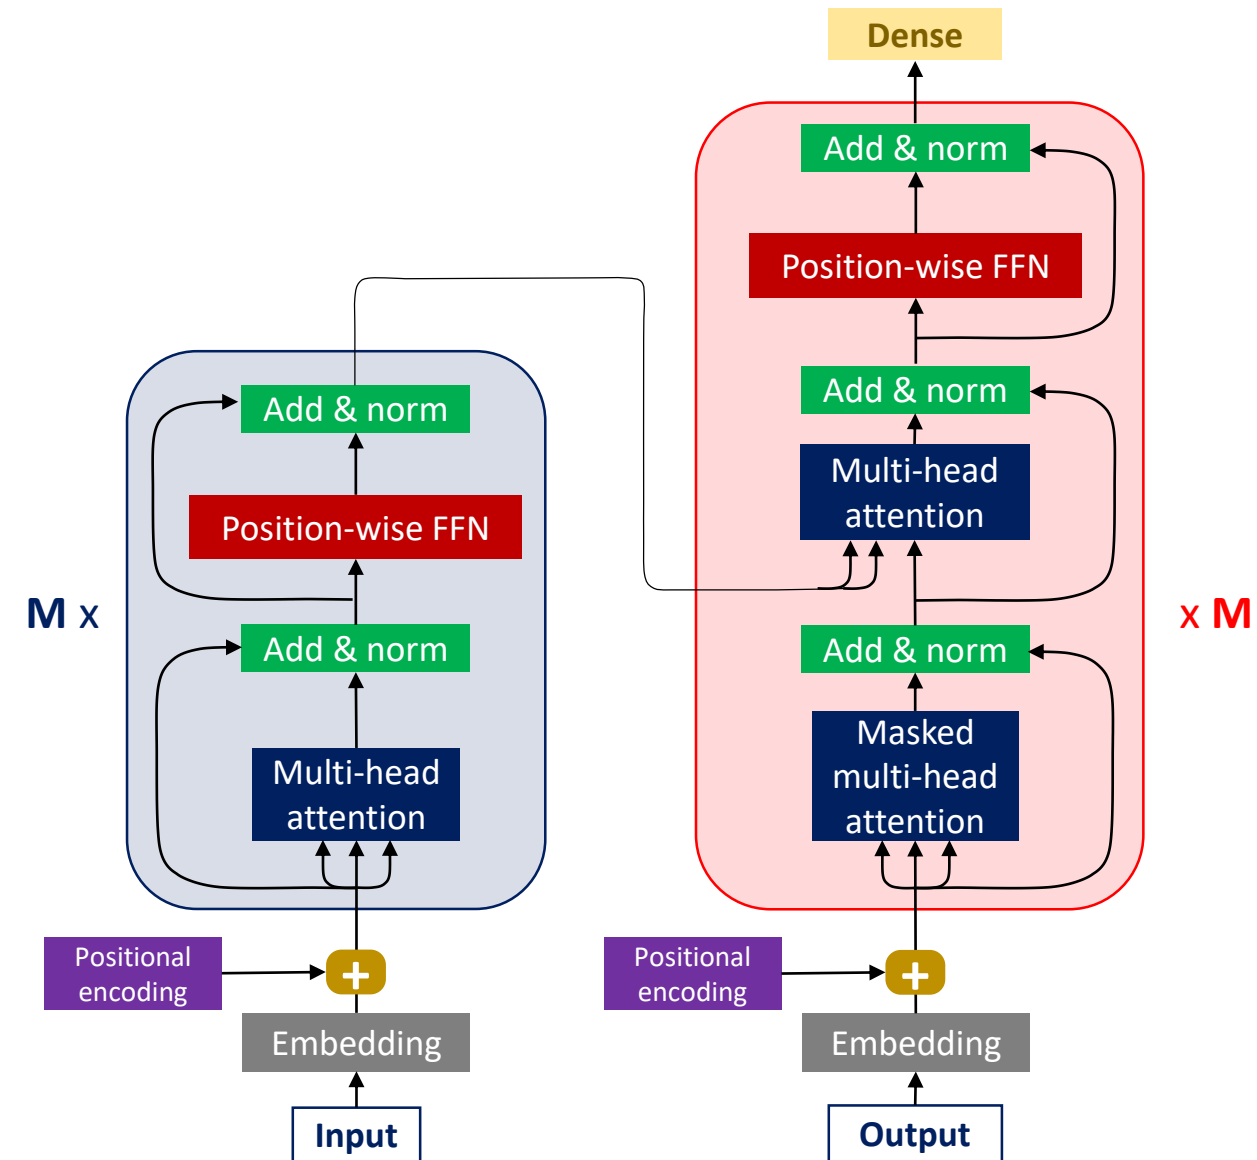


- The figure shows the basic unit of a transformer.
- Embeddings for both inputs and outputs.
- **Encoder** major components:
  - Self-attention mechanism
  - Feed-forward neural network
- **Decoder** major components:
  - Self-attention mechanism
  - Attention mechanism over encodings
  - Feed-forward neural network

# Animation

Source: <https://ai.googleblog.com/2017/08/transformer-novel-neural-network.html>


# Multi-head attention





# Multi-head attention

- In multi-head attention layer there are multiple number (say  $M$ ) of parallel self-attention layers. Each of them is called a head.
  - In a self-attention layer the query, key and value vectors are the same (the input).
- The dense layers before attention layer project the queries, keys and values into different dimensions.
  - Suppose the query, key and value vectors are projected into vectors of length  $p_q$ ,  $p_k$  and  $p_v$  respectively.
- Suppose the dimension of query is  $l_q$ , key is  $l_k$  and value is  $l_v$ .



# Multi-head attention

- Then the parameters associated with the dense layers prior to the  $m$ th attention layer are
  - $\mathbf{W}_q^m \in \mathbb{R}^{p_q \times l_q}$  for the query dense layer
  - $\mathbf{W}_k^m \in \mathbb{R}^{p_k \times l_k}$  for the key dense layer
  - $\mathbf{W}_v^m \in \mathbb{R}^{p_v \times l_v}$  for the value dense layer

- The projections for the  $m$ th head are:

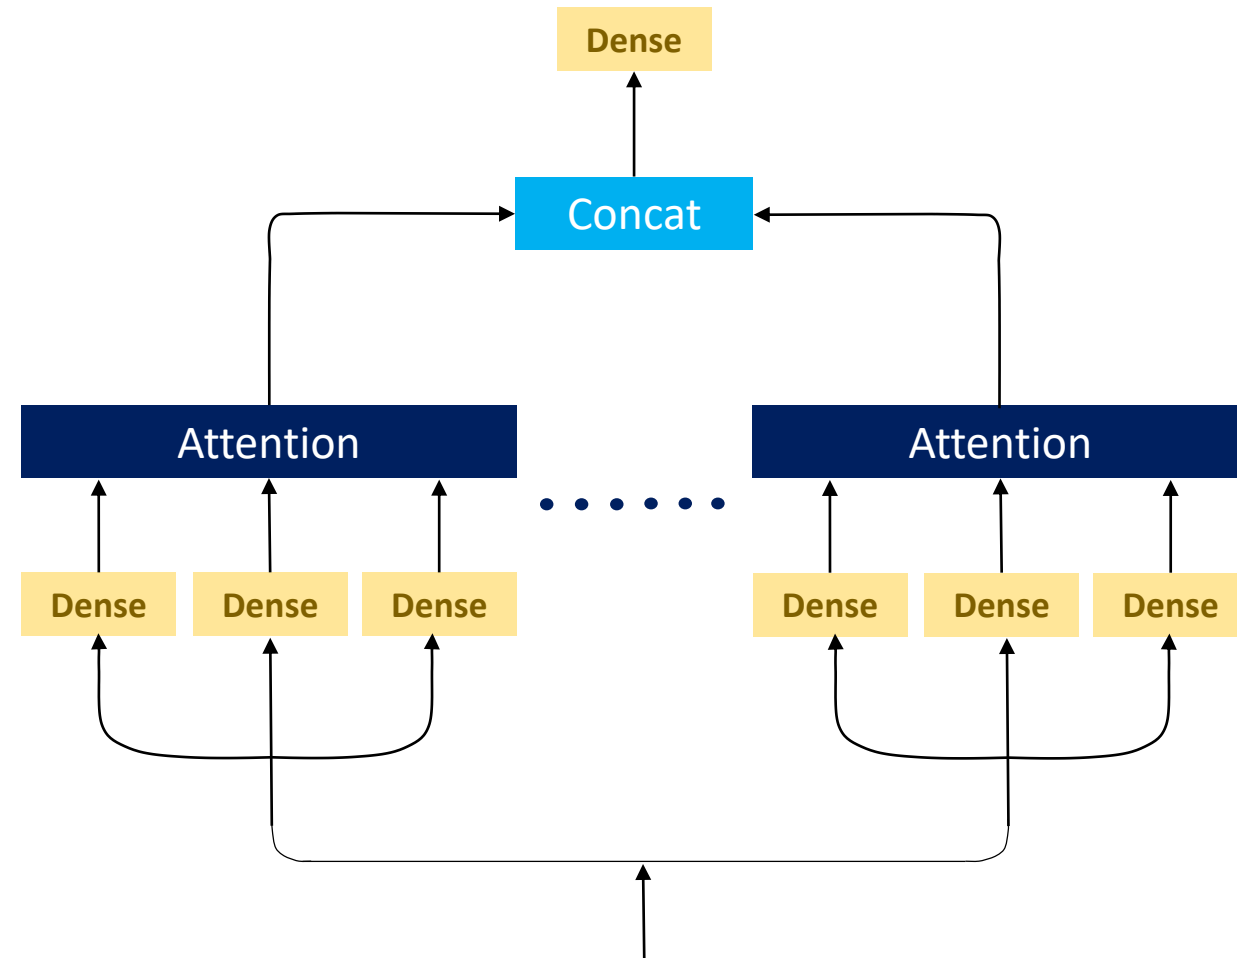
$$\mathbf{q}^{m'} = \mathbf{W}_q^m \mathbf{q}$$

$$\mathbf{k}^{m'} = \mathbf{W}_k^m \mathbf{k}$$

$$\mathbf{v}^{m'} = \mathbf{W}_v^m \mathbf{v}$$

- The output of the  $m$ th head is given as

$$\mathbf{o}^{(m)} = A(\mathbf{q}^{m'}, \mathbf{k}^{m'}, \mathbf{v}^{m'})$$

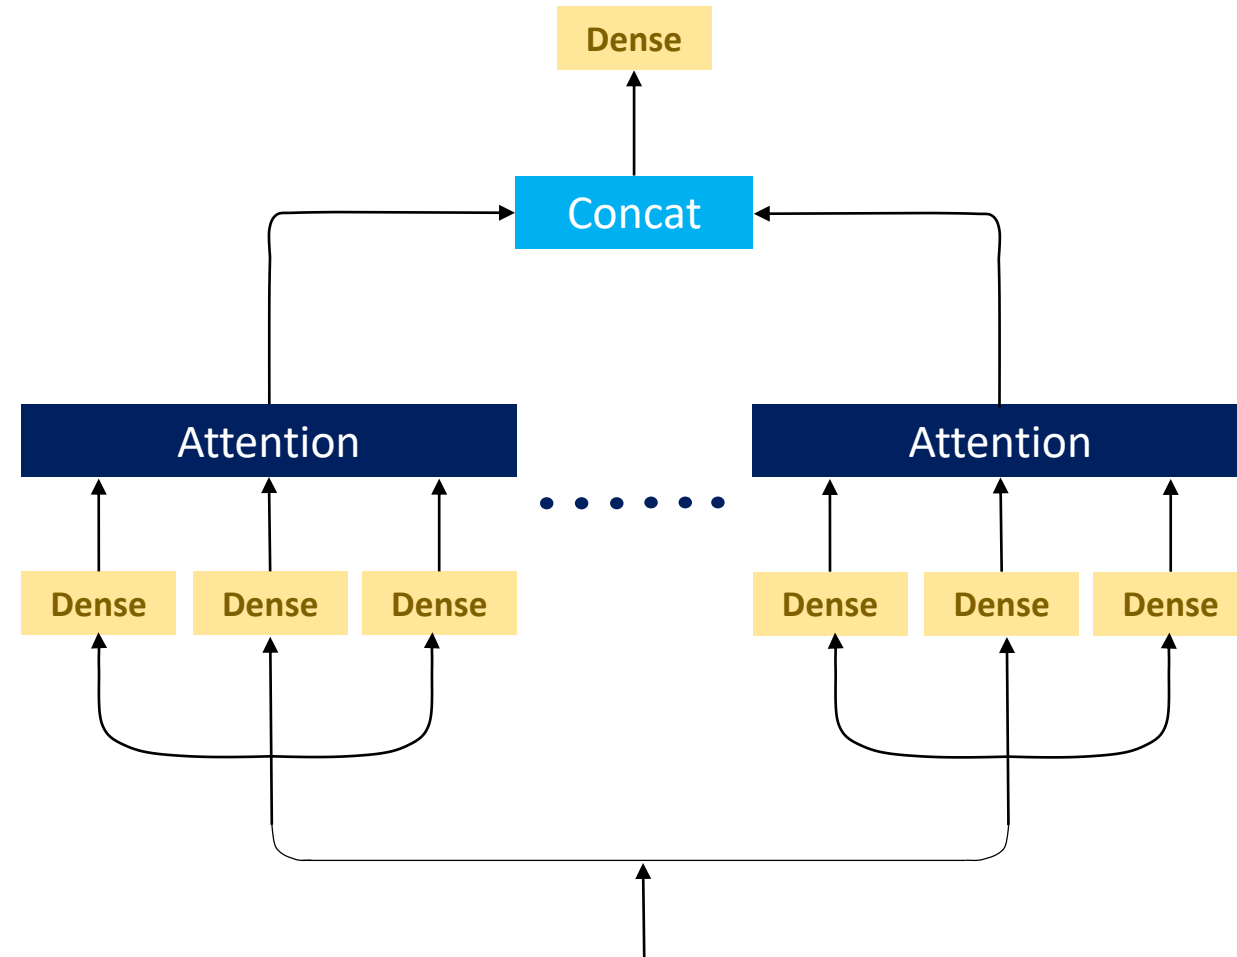


# Multi-head attention

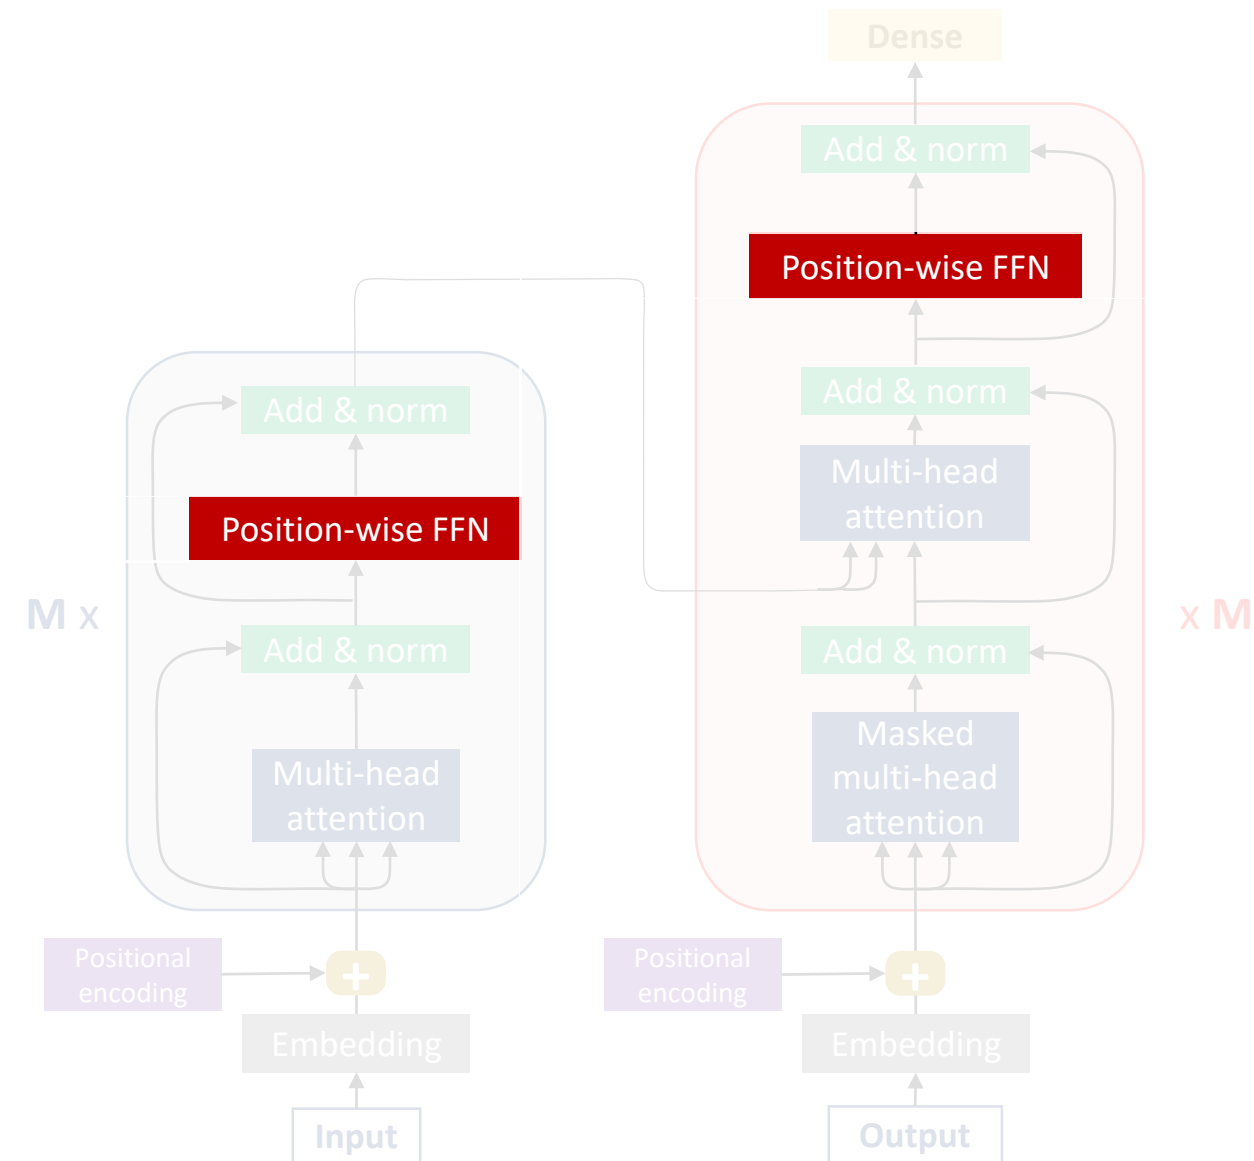
- Finally, the output of the multi-head attention

$$\mathbf{o} = \mathbf{W}_f \begin{bmatrix} \mathbf{o}^{(1)} \\ \mathbf{o}^{(2)} \\ \cdot \\ \cdot \\ \mathbf{o}^{(M)} \end{bmatrix}$$

where  $\mathbf{W}_f$  is the weight matrix corresponding to the final dense layer.


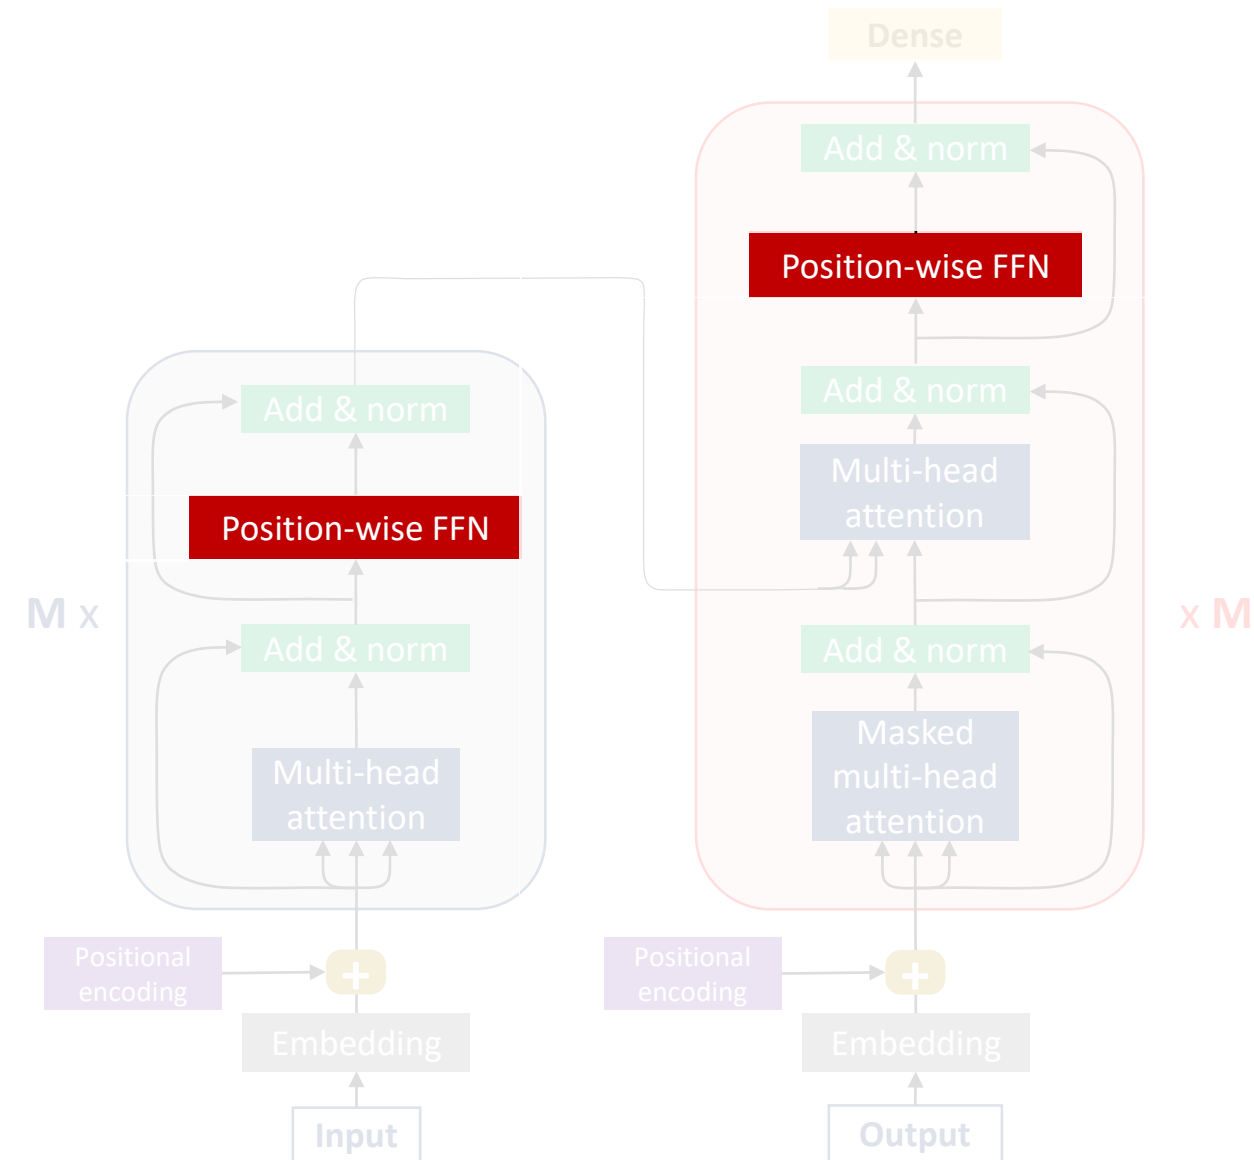


# Position-wise FFN



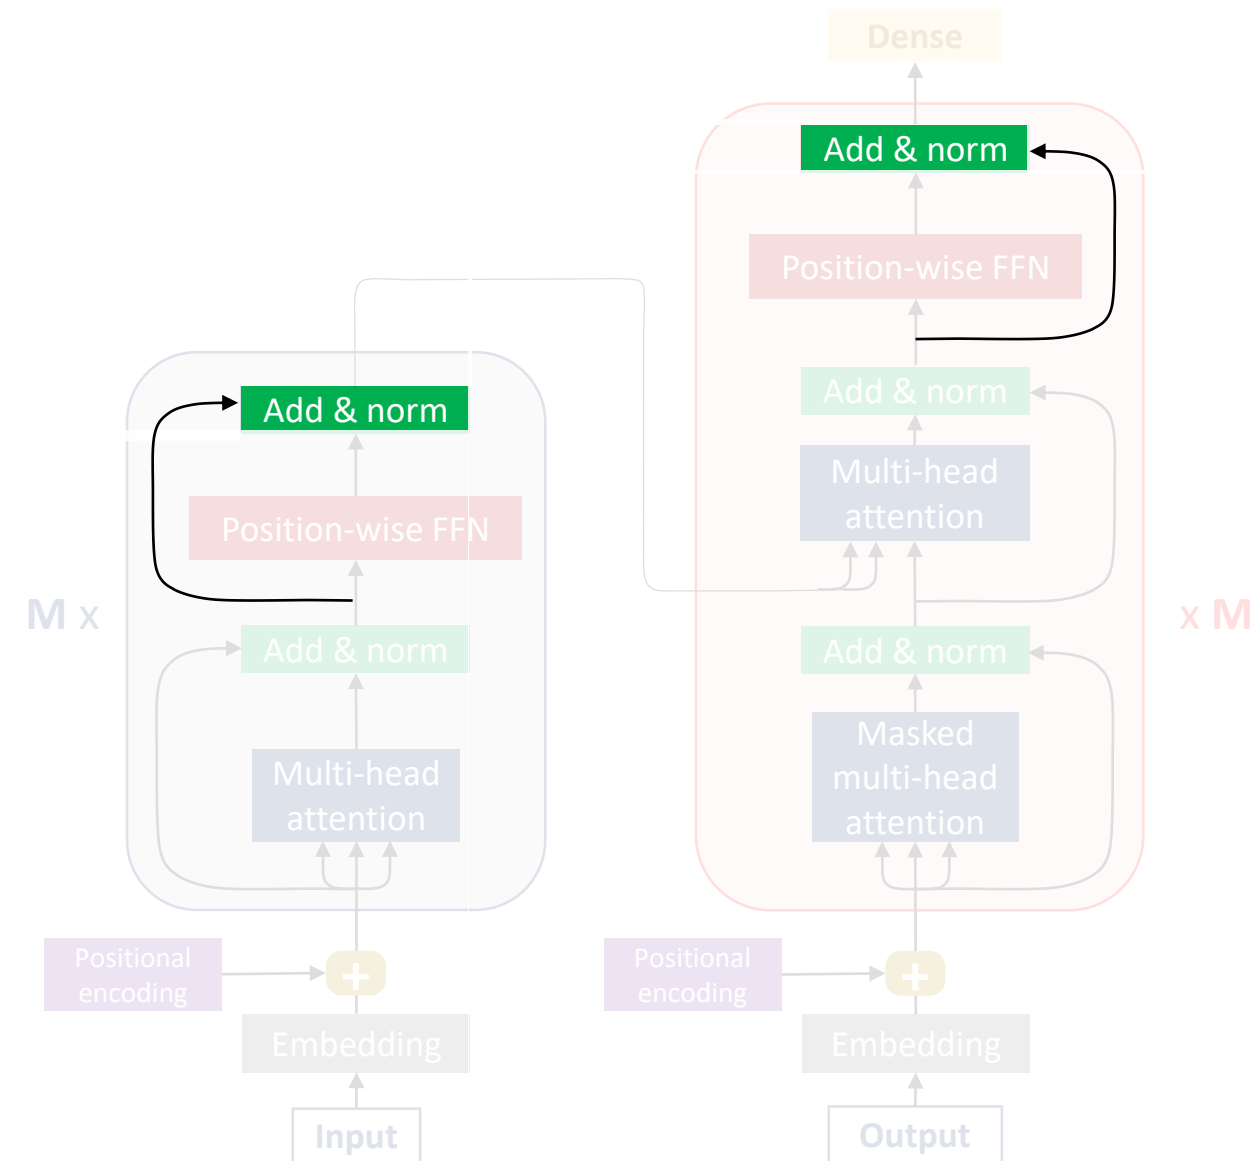
- Receives input with shape:  
(batch size, sequence length, feature dimension)
- Reshapes the input to  
(batch size  $\times$  sequence length, feature dimension)
- Pass the reshaped inputs through the two dense layers.
- Reshape back to 3D.
- Equivalent to application of two  $(1 \times 1)$  convolutions.

# Position-wise FFN



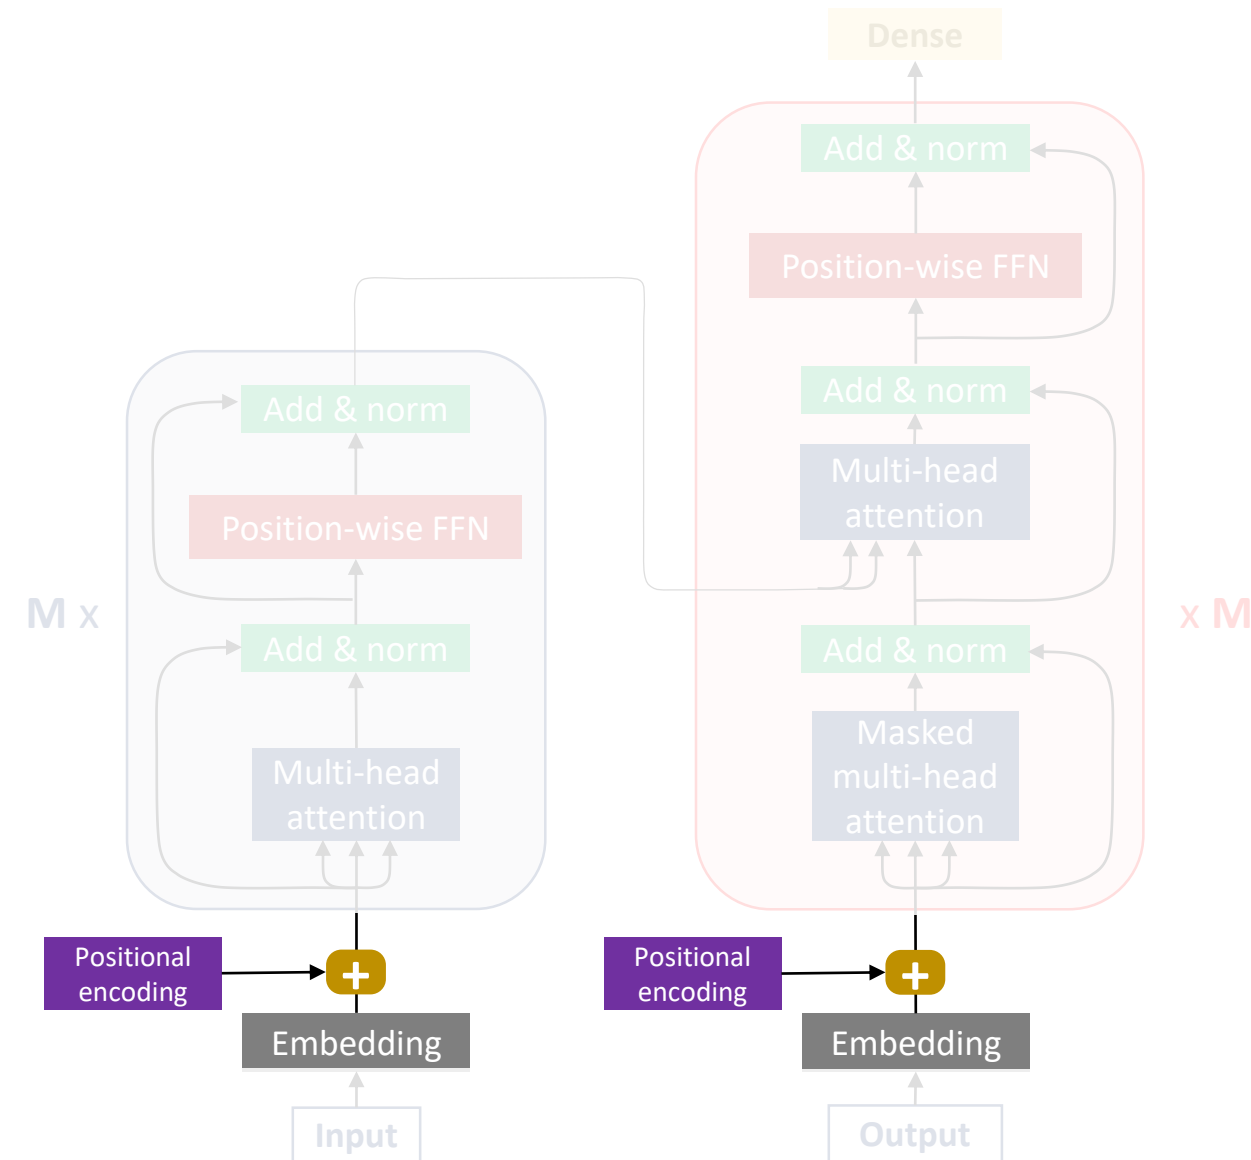


# Add & norm



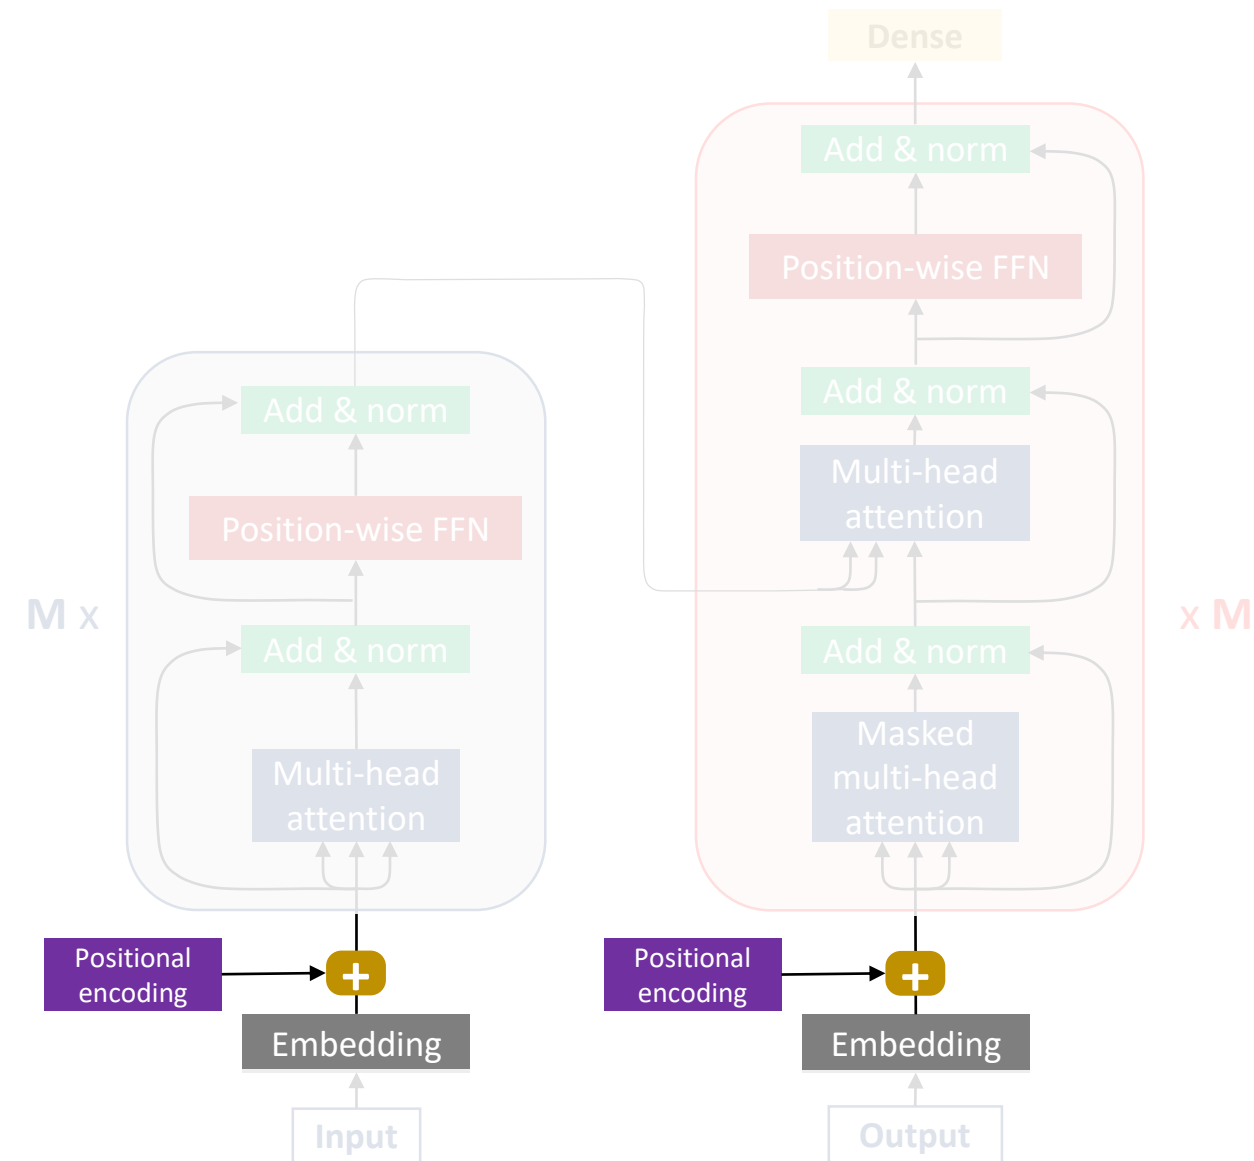
- Residual connection is employed around each of the sublayer.
- Performs layer normalization
  - Mean and variance for layer normalization are calculated w.r.t. the feature dimension.

# Positional encoding



- Positional encoding enables the transformer to **capture** sequential information.
  - The attention and the FFN layers do not retain any sequential information.
- Suppose  $\mathbf{z} \in \mathbb{R}^{s \times d_e}$  is the embedding of a given example  $\mathbf{x}$ .
  - $s$  is the length of the sequence
  - $d_e$  is the size of the embedding
- Positional encoding generates a matrix  $\mathbf{P}$  which encodes a certain form of sequential information
  - The no. of rows of  $\mathbf{P}$  is equal to sequence length.
  - The no. of columns of  $\mathbf{P}$  is equal to the embedding size.


# Positional encoding



- The matrix  $\mathbf{P}$  is defined as


$$\mathbf{P}(i, 2j) = \sin\left(\frac{i}{10000^{2j/d}}\right)$$

$$\mathbf{P}(i, 2j + 1) = \cos\left(\frac{i}{10000^{2j/d}}\right)$$



- The output of the positional encoding layer is  $\mathbf{P} + \mathbf{z}$


# Positional encoding



- The matrix  $\mathbf{P}$  is defined as

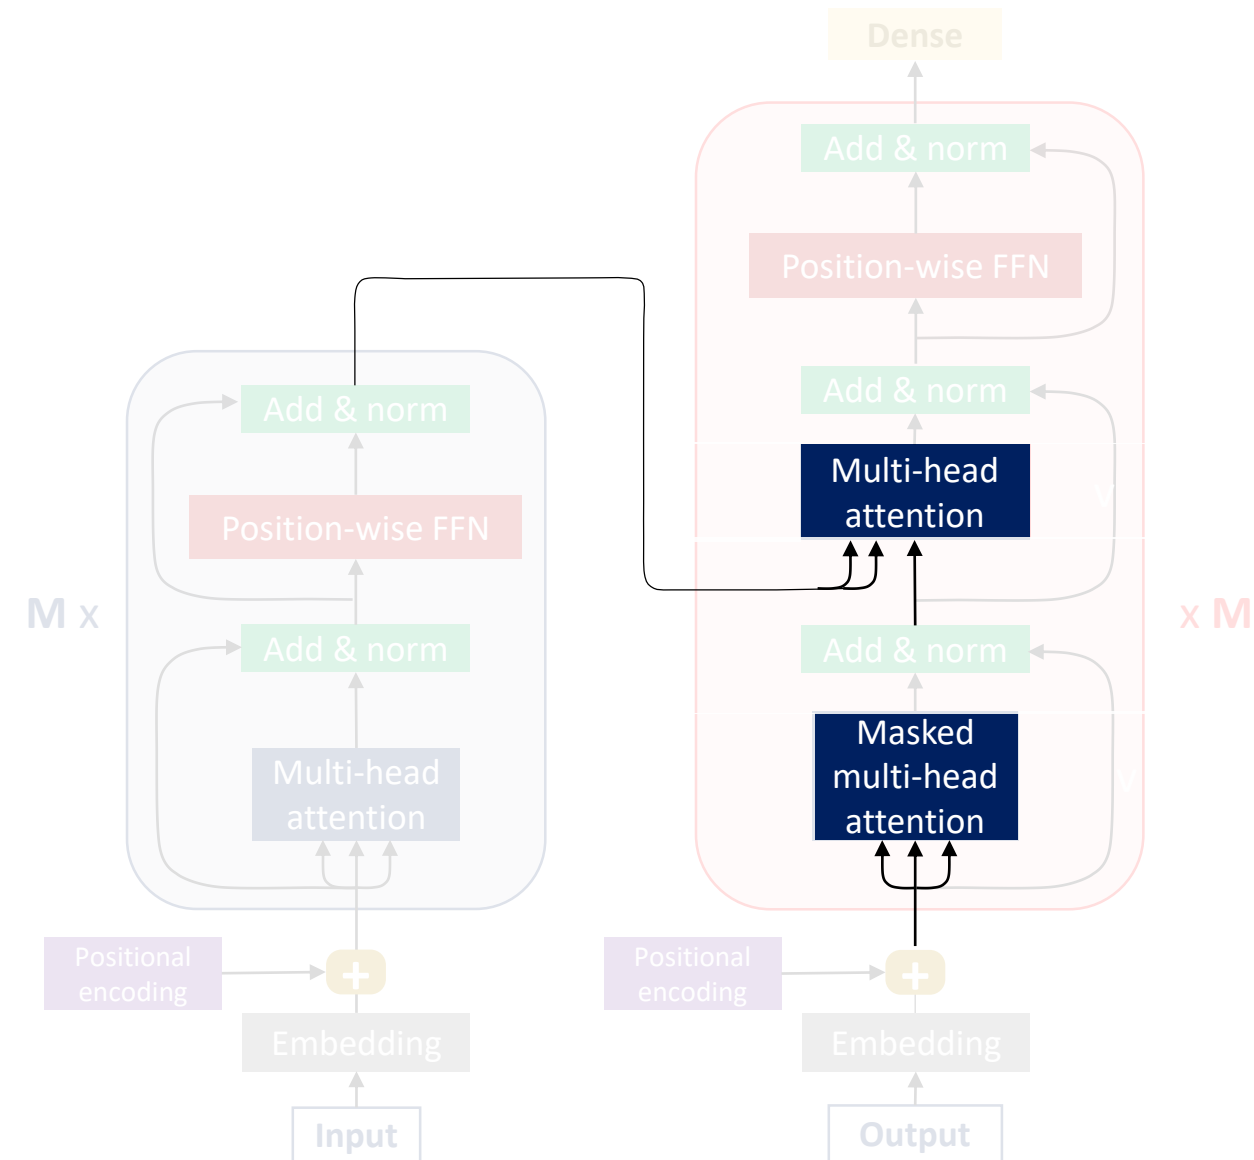
$$\mathbf{P}(i, 2j) = \sin\left(\frac{i}{10000^{2j/d}}\right)$$

$$\mathbf{P}(i, 2j + 1) = \cos\left(\frac{i}{10000^{2j/d}}\right)$$



- The output of the positional encoding layer is  $\mathbf{P} + \mathbf{z}$

# Multi-head attention layers in Decoder



- **Masked attention layer:**  
This self-attention mechanism attends to all positions in the decoder up to and including that position.
- **Encoder-Decoder attention layer:**
  - Queries come from the previous decoder layer
  - Keys and Values correspond to encoder outputs



# Flow of information

- Let us look at what part of a sentence a transformer network is attending to when processing/translating a given word.
- Consider the two sentence given below:

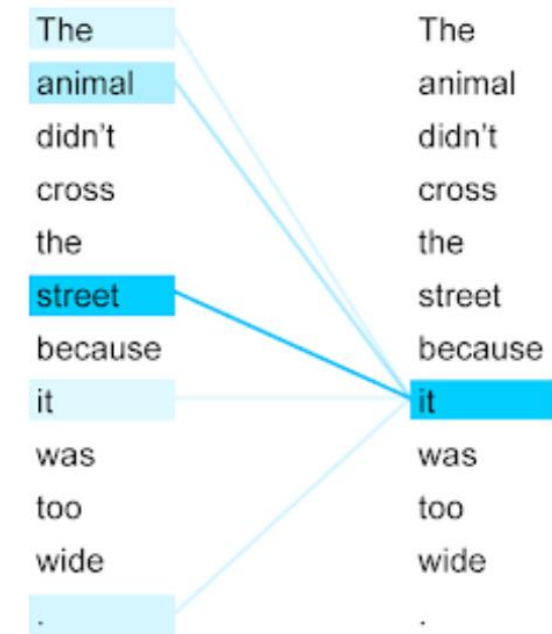
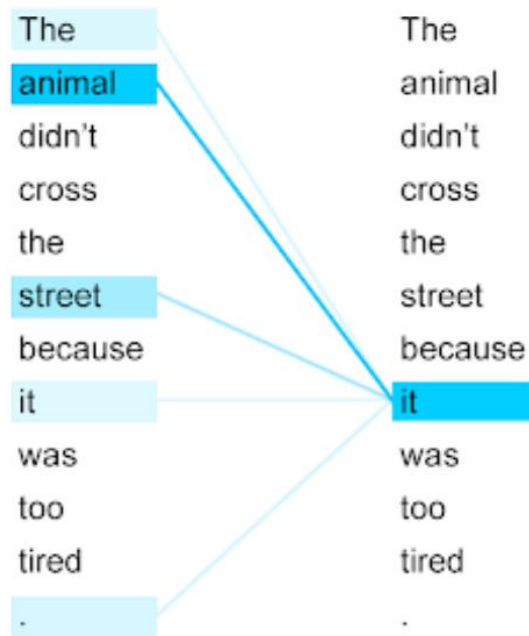
The animal didn't cross the street because it was too tired.

– Here “it” refers to the animal

The animal didn't cross the street because it was too wide.


– Here “it” refers to the street

- When translating these sentences to French, the translation for “it” depends on the gender of the noun that is referred to.

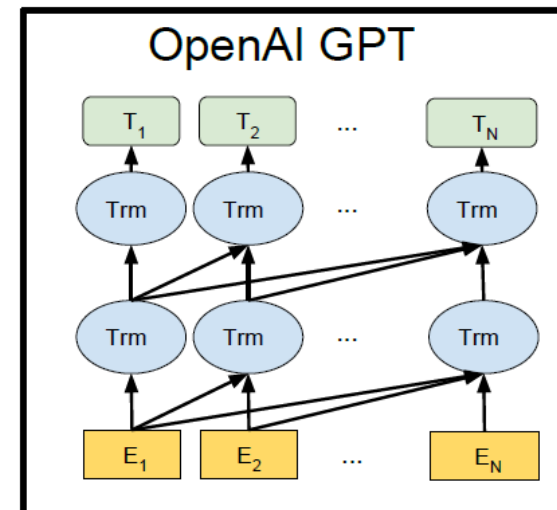


# BERT

# Earlier models

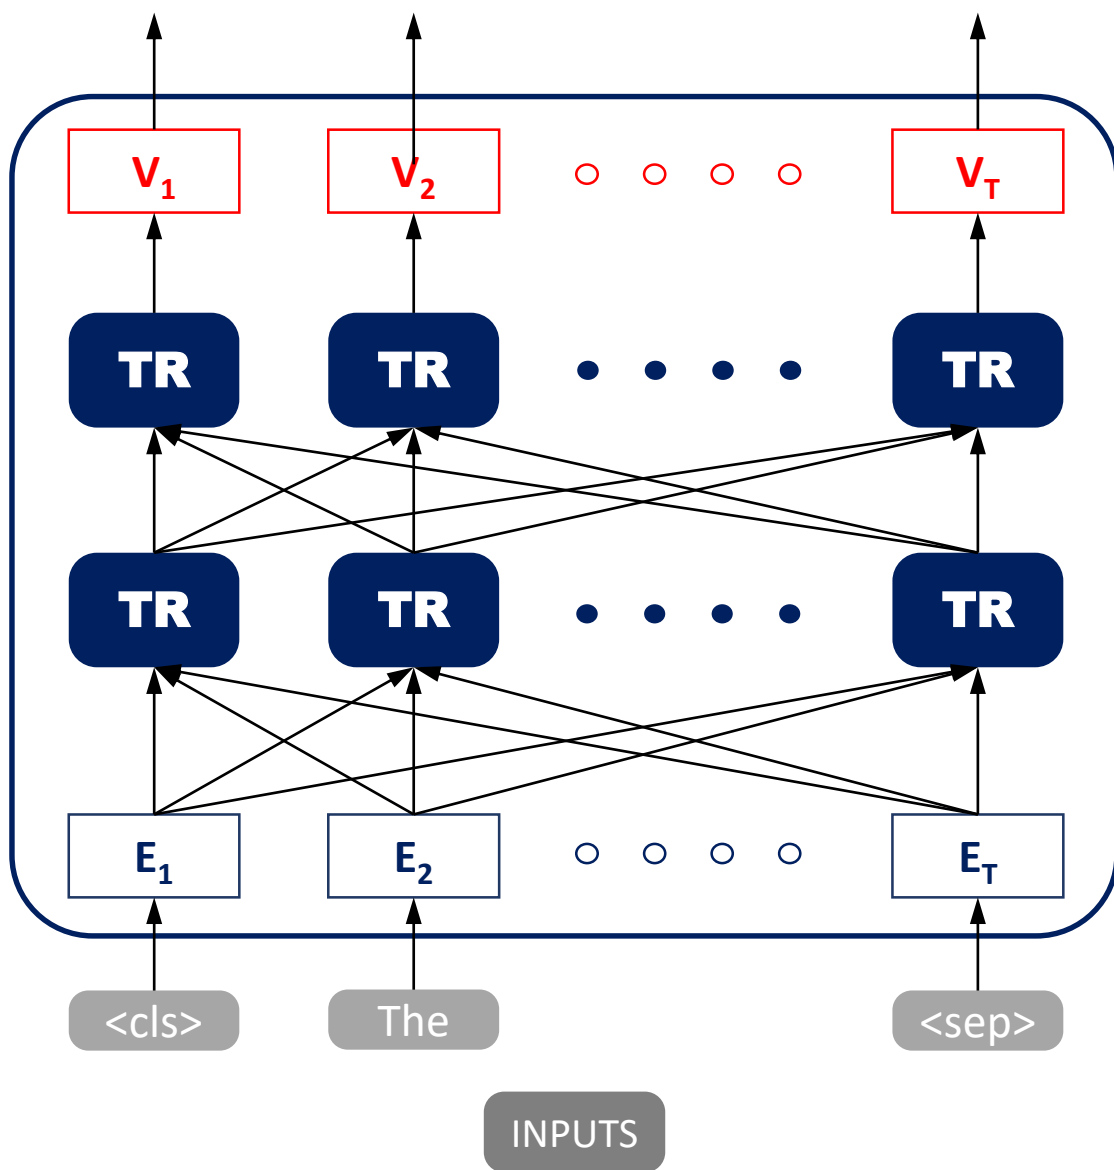


- Two language models:
  1. Left-to-right language model
  2. Right-to-left language model
- The contextual representation of each token is taken to be the concatenation of the representations of the two language models
- Both the model are trained independently



- Transformer based model
- Unidirectional: left-to-right architecture
  - Does not incorporate context from both directions

# BERT



- Bidirectional Encoder Representations from Transformers
- Architecture comprise multiple layers of transformer encoders
- Motivation: Need pre-trained language models that can achieve state-of-the-art performance.
- A big shortcoming of standard language models is **unidirectionality**.
  - Such limitations are sub-optimal for sentence level tasks.
- In many NLP tasks, e.g. question-answering, it is important the context from both sides are incorporated.

# Input representation


Net  
Embeddings



Positional  
Embeddings



Segment  
Embeddings



Token  
Embeddings




INPUTS





# Pre-training BERT: Task 1



- General procedure:
  - Some percentage of the input tokens are masked at random.
  - Final hidden vectors corresponding to the mask tokens are fed into an output softmax over the vocabulary.

# Pre-training BERT: Task 1

- The training data generator chooses 15% of tokens at random for prediction.
- If the  $t$ -th token is chosen, then it is replaced with
  1. the [MASK] token 80% of the time



2. a random token 10% of the time



3. unchanged token 10% of the time




- Finally, the vector  $\mathbf{v}_t$  is used to predict the original token.

# Pre-training BERT: Task 2

- The purpose is to make the model understand sentence relationships.
- Idea: Pre-train a binarized next-sentence prediction task.
- The training data generator chooses two sentences  $\mathcal{S}_A$  and  $\mathcal{S}_B$ , such that
  - 50% of the time  $\mathcal{S}_B$  is the actual sentence following  $\mathcal{S}_A$ . This case is labelled as *IsNext*.



- 50% of the time  $\mathcal{S}_B$  is a random sentence from the corpus. This case is labelled as *NotNext*.



# GLUE Benchmark

- Multi-Genre Natural Language Inference (MNLI): For a pair of sentences, predict if the 2nd sentence neutral w.r.t. 1st sentence.
- Quora Question Pairs (QQP): Binary classification task to determine if two questions are semantically different.
- Question Natural Language Inference (QNLI): Binary classification task where positive example are Q-A pairs with the correct answer and negative examples Q-A pairs which do not contain the answer.
- Stanford Sentiment Treebank (SST-2): Binary sentiment classification task with sentences from movie reviews.
- Corpus of Linguistic Acceptability (CoLA): Binary classification task to predict if a given sentence is linguistically acceptable.
- Semantic Textual Similarity Benchmark (STS-B):: Indicate the (semantic) similarity of 2 sentences on a scale of 1 to 5.

# Pre-training → Fine-Tuning

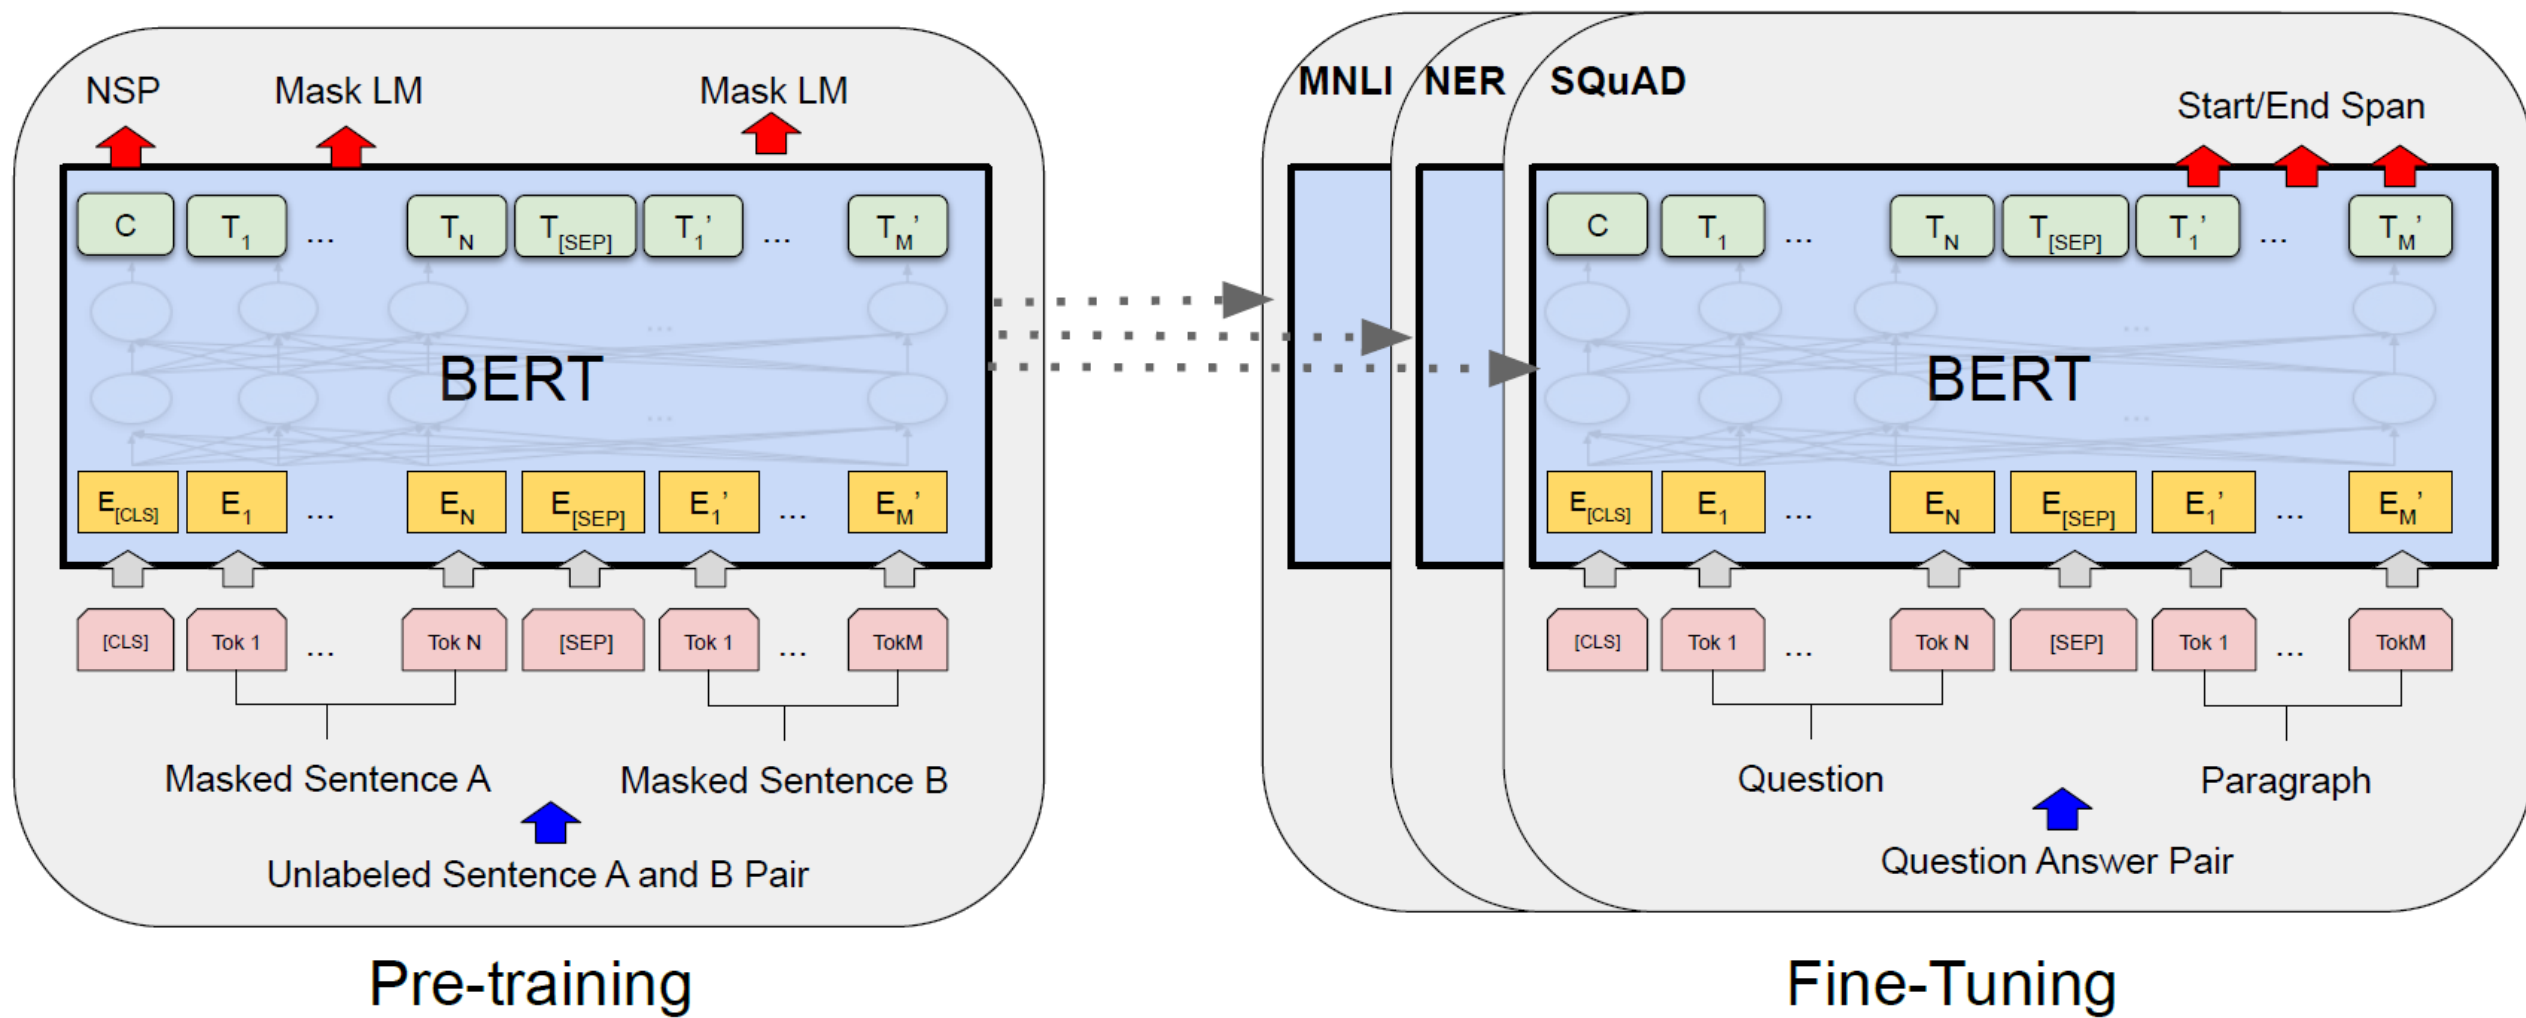


Figure source: Devlin *et. al.* BERT: Pre-training Deep Bidirectional Transformers for Language Understanding, Proceedings of NAACL-HLT 2019