

Data Frames

- Data frames are used to store tabular data in R. Hadley Wickham's package `dplyr`³⁵ has an optimized set of functions designed to work efficiently with data frames.
- Data frames are represented as a special type of list where every element of the list has to have the same length. Each element of the list can be thought of as a column and the length of each element of the list is the number of rows.
- Unlike matrices, data frames can store different classes of objects in each column. Matrices must have every element be the same class (e.g. all integers or all numeric).

Data Frames

- In addition to column names, indicating the names of the variables or predictors, data frames have a special attribute called `row.names` which indicate information about each row of the data frame.
- Data frames are usually created by reading in a dataset using the `read.table()` or `read.csv()`. However, data frames can also be created explicitly with the `data.frame()` function.

Data Frames: Construction

- `x <- data.frame(foo = 1:4, bar = c(T, T, F, F))`

- `x`

	foo	bar
1	1	TRUE
2	2	TRUE
3	3	FALSE
4	4	FALSE

- `nrow(x)`

`[1] 4`

- `ncol(x)`

`[1] 2`

Accessing Data Frame

- `vector1 <- c(10,30,40,15); vector1`
- `vector2 <- vector1 * 2`
- `df1 <- data.frame(vector1,vector2); df1`
- `vector3 <- c(1,2,3)`
- `vector3 <- c(1,2,3,4)`
- `df1 <- data.frame(vector1,vector2,vector3); df1` # gives error
- `df1[1,2]`
- `df1[,2]`

Accessing Data Frame

- `df1$vector1` #similar to the use of `$` in the case of list, because `df` is a list
- `df1[df1$vector1<20,2]`
- `df1`
- `df1$vector1<20`
- `df1[1:3,2]`
- `is.list(df1$vector1[1:2])`
- `is.vector(df1$vector1[1:2])`
- `str(df1)`

Removing NAs in Data Frames

- `head(airquality)`
- `good <- complete.cases(airquality)`
- `head(airquality[good,])`
- `x <- airquality[, -1]` # x is a regression design matrix
- `y <- airquality[, 1]` # y is the corresponding response
- `stopifnot(complete.cases(y) != is.na(y))`
- `ok <- complete.cases(x, y)`
- `sum(!ok)` # how many are not "ok" ?
- `x <- x[ok,]`
- `y <- y[ok]`

Reading and Writing Data

There are a few principal functions reading data into R.

- `read.table`, `read.csv` , for reading tabular data
- `readLines` , for reading lines of a text file

Reading and Writing Data

There are analogous functions for writing data to files

- `write.table` , for writing tabular data to text files (i.e. CSV) or connections
- `writeln` , for writing character data line-by-line to a file or connection

read.table – important parameters

- `file` , the name of a file, or a connection
- `header` , logical indicating if the file has a header line
- `sep` , a string indicating how the columns are separated
- `colClasses` , a character vector indicating the class of each column in the dataset
- `na.strings` – the set of strings to be considered as NA

Reading Data

- `tab1 <- read.table('house_copy.txt')`
- `tab1`
- `getwd()`
- `str(tab1)`
- `tab1 <- read.table(file='house_copy.txt',header = T)`
- `tab1 <- read.table(file='house_copy.txt',header = T, sep = '', colClasses = c('double','double'))`
- `tab1 <- read.table(file='house_copy.txt',header = T, sep = '', colClasses = c('double','double'), na.strings = c('220.0'))`
- `tab1 <- read.table("home/user/house.txt")` #using relative path notation
- #please experiment with relative/absolute path in Windows - use `getwd()` for hints
- #experiment with `read.csv` and csv files (comma-separated files like from excel)
- `vector1 <- c(10,30,40,15); vector1`
- `vector1`
- `writeLines(as.character(vector1), 'vector1.txt')`
- `getwd()`
- `write.table(tab1,file = 'area_sales.txt',sep=',')`
- `write.table(tab1,file = 'area_sales.txt',sep=',',quote = F,row.names = F)`

Exercise

- Experiment with variations in write.table function (command) to get the data in the original format as shown below:

area,sale.price

694,192

905,215

802,215

1366,274

716,112.7

963,185

821,212

714,220

1018,276

887,260

790,221.5

696,255

771,260

1006,293

1191,375

- Experiment reading CSV files using read.csv function and explore usage of its various parameters (arguments)

Testing and Coercion – compilation

Table 2.3. Functions for testing (`is`) the attributes of different categories of object (arrays, lists, etc.) and for coercing (`as`) the attributes of an object into a specified form. Neither operation changes the attributes of the object unless you overwrite its name.

Type	Testing	Coercing
Array	<code>is.array</code>	<code>as.array</code>
Character	<code>is.character</code>	<code>as.character</code>
Complex	<code>is.complex</code>	<code>as.complex</code>
Dataframe	<code>is.data.frame</code>	<code>as.data.frame</code>
Double	<code>is.double</code>	<code>as.double</code>
Factor	<code>is.factor</code>	<code>as.factor</code>
List	<code>is.list</code>	<code>as.list</code>
Logical	<code>is.logical</code>	<code>as.logical</code>
Matrix	<code>is.matrix</code>	<code>as.matrix</code>
Numeric	<code>is.numeric</code>	<code>as.numeric</code>
Raw	<code>is.raw</code>	<code>as.raw</code>
Time series (ts)	<code>is.ts</code>	<code>as.ts</code>
Vector	<code>is.vector</code>	<code>as.vector</code>

Table 2.4. Vector functions used in R.

Operation	Meaning
<code>max(x)</code>	maximum value in x
<code>min(x)</code>	minimum value in x
<code>sum(x)</code>	total of all the values in x
<code>mean(x)</code>	arithmetic average of the values in x
<code>median(x)</code>	median value in x
<code>range(x)</code>	vector of <code>min(x)</code> and <code>max(x)</code>
<code>var(x)</code>	sample variance of x
<code>cor(x, y)</code>	correlation between vectors x and y
<code>sort(x)</code>	a sorted version of x
<code>rank(x)</code>	vector of the ranks of the values in x
<code>order(x)</code>	an integer vector containing the permutation to sort x into ascending order
<code>quantile(x)</code>	vector containing the minimum, lower quartile, median, upper quartile, and maximum of x
<code>cumsum(x)</code>	vector containing the sum of all of the elements up to that point
<code>cumprod(x)</code>	vector containing the product of all of the elements up to that point
<code>cummax(x)</code>	vector of non-decreasing numbers which are the cumulative maxima of the values in x up to that point
<code>cummin(x)</code>	vector of non-increasing numbers which are the cumulative minima of the values in x up to that point
<code>pmax(x, y, z)</code>	vector, of length equal to the longest of x , y or z , containing the maximum of x , y or z for the i th position in each
<code>pmin(x, y, z)</code>	vector, of length equal to the longest of x , y or z , containing the minimum of x , y or z for the i th position in each
<code>colMeans(x)</code>	column means of dataframe or matrix x
<code>colSums(x)</code>	column totals of dataframe or matrix x
<code>rowMeans(x)</code>	row means of dataframe or matrix x
<code>rowSums(x)</code>	row totals of dataframe or matrix x